

# Data Mining 2017

## Text Classification

### Naive Bayes

Ad Feelders

Universiteit Utrecht

October 18, 2017

- Text Mining is data mining applied to text data.
- Often uses well-known data mining algorithms.
- Text data requires substantial pre-processing.
- This typically results in a large number of attributes (for example, the size of the dictionary).

# Text Classification

- Predict the class(es) of text documents.
- Can be single-label or multi-label.
- Multi-label classification is often performed by building multiple binary classifiers (one for each possible class).
- Examples of text classification:
  - topics of news articles,
  - spam/no spam for e-mail messages,
  - sentiment analysis (e.g. positive/negative review),
  - opinion spam (e.g. fake reviews),
  - music genre from song lyrics

# Is this Rap, Blues, Metal, Country or Pop?

Blasting our way through the boundaries of Hell  
No one can stop us tonight  
We take on the world with hatred inside  
Mayhem the reason we fight  
Surviving the slaughters and killing we've lost  
Then we return from the dead  
Attacking once more now with twice as much strength  
We conquer then move on ahead

[Chorus:]

Evil  
My words defy  
Evil  
Has no disguise  
Evil  
Will take your soul  
Evil  
My wrath unfolds

Satan our master in evil mayhem  
Guides us with every first step  
Our axes are growing with power and fury  
Soon there'll be nothingness left  
Midnight has come and the leathers strapped on  
Evil is at our command  
We clash with God's angel and conquer new souls  
Consuming all that we can

## Probabilistic Classifier

A probabilistic classifier assigns a probability to each class. In case a class prediction is required we typically predict the class with highest probability:

$$\hat{c} = \arg \max_{c \in C} P(c | d) = \arg \max_{c \in C} \frac{P(d | c)P(c)}{P(d)}$$

where  $d$  is a document, and  $C$  is the set of all possible class labels.

Since  $P(d) = \sum_{c \in C} P(c, d)$  is the same for all classes, we can ignore the denominator:

$$\hat{c} = \arg \max_{c \in C} P(c | d) = \arg \max_{c \in C} P(d | c)P(c)$$

# Naive Bayes

Represent document as set of features:

$$\hat{c} = \arg \max_{c \in C} P(c \mid d) = \arg \max_{c \in C} P(x_1, \dots, x_m \mid c)P(c)$$

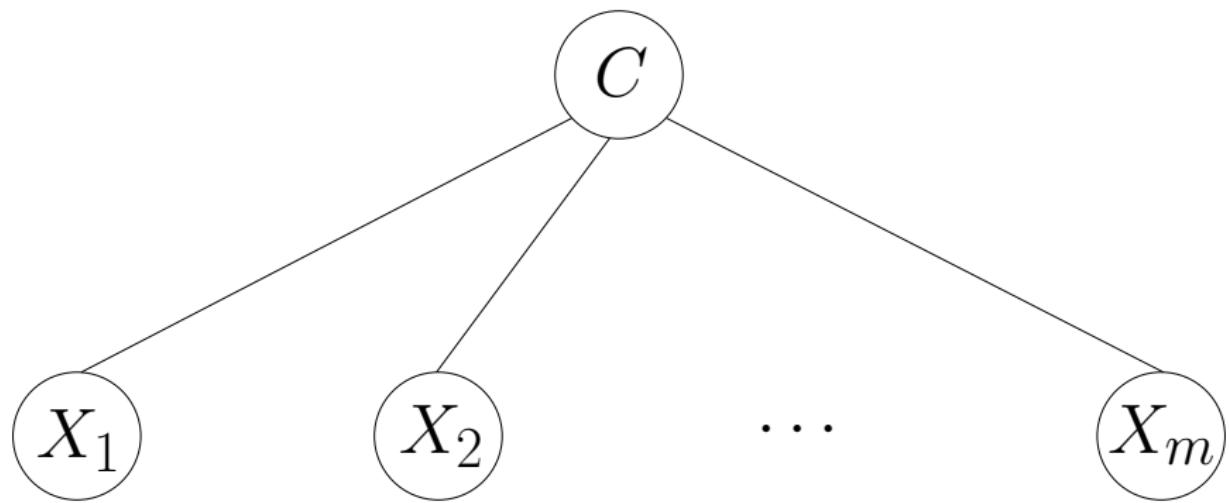
Naive Bayes assumption:

$$P(x_1, \dots, x_m \mid c) = P(x_1 \mid c)P(x_2 \mid c) \cdot \dots \cdot P(x_m \mid c)$$

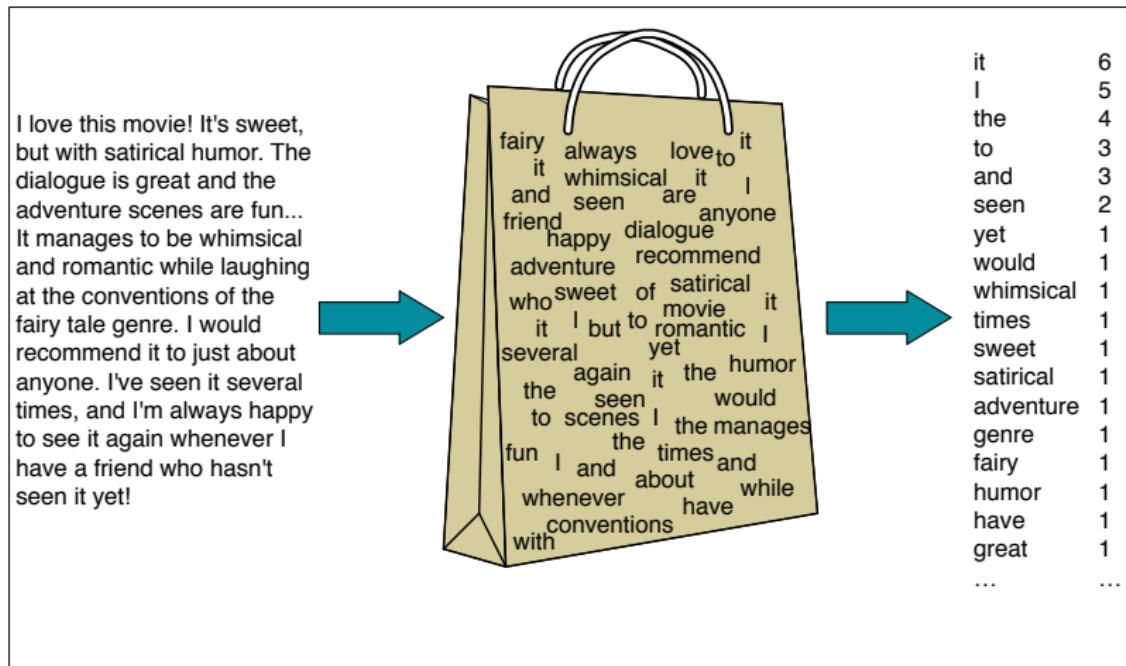
The features are independent within each class (avoiding the curse of dimensionality).

$$c_{\text{NB}} = \arg \max_{c \in C} P(c) \prod_{i=1}^m P(x_i \mid c)$$

# Independence Graph of Naive Bayes



# Bag Of Words Representation of a Document



**Figure 6.1** Intuition of the multinomial naive Bayes classifier applied to a movie review. The position of the words is ignored (*the bag of words* assumption) and we make use of the frequency of each word.

# Multinomial Naive Bayes for Text

Represent document  $d$  as a sequence of words:  $d = \langle w_1, w_2, \dots, w_n \rangle$ .

$$c_{\text{NB}} = \arg \max_{c \in C} P(c) \prod_{k=1}^n P(w_k | c)$$

Notice that  $P(w | c)$  is independent of word position or word order, so  $d$  is truly represented as a bag-of-words. Taking the log we obtain:

$$c_{\text{NB}} = \arg \max_{c \in C} \log P(c) + \sum_{k=1}^n \log P(w_k | c)$$

By the way, why is it allowed to take the log?

# Multinomial Naive Bayes for Text

Consider the text (perhaps after some pre-processing)

catch as catch can

We have  $d = \langle \text{catch}, \text{as}, \text{catch}, \text{can} \rangle$ , with  $w_1 = \text{catch}$ ,  $w_2 = \text{as}$ ,  $w_3 = \text{catch}$ , and  $w_4 = \text{can}$ . Suppose we have two classes, say  $C = \{+, -\}$ , then for this document:

$$\begin{aligned}c_{\text{NB}} &= \arg \max_{c \in \{+, -\}} \log P(c) + \log P(\text{catch} | c) + \log P(\text{as} | c) \\&\quad + \log P(\text{catch} | c) + \log P(\text{can} | c) \\&= \arg \max_{c \in \{+, -\}} \log P(c) + 2 \log P(\text{catch} | c) + \log P(\text{as} | c) \\&\quad + \log P(\text{can} | c)\end{aligned}$$

# Training Multinomial Naive Bayes

Class priors:

$$\hat{P}(c) = \frac{N_c}{N_{doc}}$$

Word probabilities within each class:

$$\hat{P}(w_i | c) = \frac{\text{count}(w_i, c)}{\sum_{w_j \in V} \text{count}(w_j, c)},$$

Where  $V$  (for Vocabulary) denotes the collection of all words that occur in the training corpus (after possibly extensive pre-processing).

Verify that

$$\sum_{w_i \in V} \hat{P}(w_i | c) = 1,$$

as required.

## Training Multinomial Naive Bayes

Perform *smoothing* to avoid zero probability estimates.

Word probabilities within each class with Laplace smoothing:

$$\hat{P}(w_i | c) = \frac{\text{count}(w_i, c) + 1}{\sum_{w_j \in V} (\text{count}(w_j, c) + 1)} = \frac{\text{count}(w_i, c) + 1}{\sum_{w_j \in V} \text{count}(w_j, c) + |V|}$$

Verify that again

$$\sum_{w_i \in V} \hat{P}(w_i | c) = 1,$$

as required.

# Multinomial Naive Bayes: Training

TRAINMULTINOMIALNB( $C, D$ )

```
1   $V \leftarrow \text{EXTRACTVOCABULARY}(D)$ 
2   $N_{\text{doc}} \leftarrow \text{COUNTDOCS}(D)$ 
3  for each  $c \in C$ 
4  do  $N_c \leftarrow \text{COUNTDOCSINCLASS}(D, c)$ 
5     $\text{prior}[c] \leftarrow N_c / N_{\text{doc}}$ 
6     $\text{text}_c \leftarrow \text{CONCATENATETEXTOFAALLDOCSINCLASS}(D, c)$ 
7    for each  $w \in V$ 
8      do  $\text{count}_{cw} \leftarrow \text{COUNTWORDOCCURRENCE}(\text{text}_c, w)$ 
9      for each  $w \in V$ 
10     do  $\text{condprob}[w][c] \leftarrow \frac{\text{count}_{cw} + 1}{\sum_{w'} (\text{count}_{cw'} + 1)}$ 
11  return  $V, \text{prior}, \text{condprob}$ 
```

# Multinomial Naive Bayes: Prediction

Predict the class of a document  $d$ .

APPLYMULTINOMIALNB( $C, V, prior, condprob, d$ )

- 1     $W \leftarrow \text{EXTRACTWORDOCCURRENCESFROMDOC}(V, d)$
- 2    **for each**  $c \in C$
- 3       **do**  $score[c] \leftarrow \log prior[c]$
- 4          **for each**  $w \in W$
- 5             **do**  $score[c] += \log condprob[w][c]$
- 6    **return**  $\arg \max_{c \in C} score[c]$

## Worked Example: Movie Reviews

Cat	Documents
Training	<ul style="list-style-type: none"><li>- just plain boring</li><li>- entirely predictable and lacks energy</li><li>- no surprises and very few laughs</li><li>+ very powerful</li><li>+ the most fun film of the summer</li></ul>
Test	? predictable with no fun

# Class Prior Probabilities

Recall that:

$$\hat{P}(c) = \frac{N_c}{N_{doc}}$$

So we get:

$$\hat{P}(+) = \frac{2}{5} \quad \hat{P}(-) = \frac{3}{5}$$

# Word Conditional Probabilities

To classify the test example, we need the following probability estimates:

$$\hat{P}(\text{"predictable"} \mid -) = \frac{1+1}{14+20} = \frac{1}{17} \quad \hat{P}(\text{"predictable"} \mid +) = \frac{0+1}{9+20} = \frac{1}{29}$$

$$\hat{P}(\text{"no"} \mid -) = \frac{1+1}{14+20} = \frac{1}{17} \quad \hat{P}(\text{"no"} \mid +) = \frac{0+1}{9+20} = \frac{1}{29}$$

$$\hat{P}(\text{"fun"} \mid -) = \frac{0+1}{14+20} = \frac{1}{34} \quad \hat{P}(\text{"fun"} \mid +) = \frac{1+1}{9+20} = \frac{2}{29}$$

Classification:

$$\hat{P}(-)\hat{P}(\text{predictable no fun} \mid -) = \frac{3}{5} \times \frac{1}{17} \times \frac{1}{17} \times \frac{1}{34} = \frac{3}{49,130}$$

$$\hat{P}(+)\hat{P}(\text{predictable no fun} \mid +) = \frac{2}{5} \times \frac{1}{29} \times \frac{1}{29} \times \frac{2}{29} = \frac{4}{121,945}$$

The model predicts class *negative* for the test review.

# Violation of Naive Bayes independence assumptions

The multinomial naive Bayes model makes two kinds of independence assumptions:

- ① Conditional independence:

$$P(\langle w_1, \dots, w_n \rangle | c) = \prod_{k=1}^n P(W_k = w_k | c)$$

- ② Positional independence:  $P(W_{k_1} = w | c) = P(W_{k_2} = w | c)$

These independence assumptions do not really hold for documents written in natural language.

How can naive Bayes get away with such *heroic* assumptions?

# Why does Naive Bayes work?

- Naive Bayes can work well even though independence assumptions are *badly* violated.
- Example:

	$c_1$	$c_2$	predicted
true probability $P(c d)$	0.6	0.4	$c_1$
$\hat{P}(c) \prod \hat{P}(w_k c)$	0.00099	0.00001	
NB estimate $\hat{P}(c d)$	0.99	0.01	$c_1$

- Double counting of evidence causes underestimation (0.01) and overestimation (0.99).
- Classification is about predicting the correct class, *not* about accurate estimation.

## Naive Bayes is not so naive

- Probability estimates may be way off, but that doesn't have to hurt classification performance (much).
- Requires the estimation of relatively few parameters, which may be beneficial if you have a small training set.
- Fast, low storage requirements

# Feature Selection

The vocabulary of a training corpus may be huge, but not all words will be good class predictors.

How can we reduce the number of features?

- Feature utility measures:
  - Frequency – select the most frequent terms.
  - Mutual information – select the terms that have the highest mutual information with the class label.
  - Chi-square test of independence between term and class label.
- Sort features by utility and select top  $k$ .
- Can we miss good sets of features this way?

# Entropy and Conditional Entropy

Entropy is the average amount of information generated by observing the value of a random variable:

$$H(X) = \sum_x P(x) \log_2 \frac{1}{P(x)} = - \sum_x P(x) \log_2 P(x)$$

We can also interpret it as a measure of the uncertainty about the value of  $X$  prior to observation.

Conditional entropy:

$$H(X | Y) = \sum_{x,y} P(x,y) \log_2 \frac{1}{P(x | y)} = - \sum_{x,y} P(x,y) \log_2 P(x | y)$$

# Mutual Information

For random variables  $X$  and  $Y$ , their mutual information is given by

$$I(X; Y) = H(X) - H(X | Y) = \sum_x \sum_y P(x, y) \log_2 \frac{P(x, y)}{P(x)P(y)}$$

- Mutual information measures the reduction in uncertainty about  $X$  achieved by observing the value of  $Y$  (and vice versa).
- If  $X$  and  $Y$  are independent, then for all  $x, y$  we have  $P(x, y) = P(x)P(y)$ , so  $I(X, Y) = 0$ .
- Otherwise  $I(X; Y)$  is a positive quantity.

# Estimated Mutual Information

To estimate  $I(X; Y)$  from data we compute

$$I(X; Y) = \sum_x \sum_y \hat{P}(x, y) \log_2 \frac{\hat{P}(x, y)}{\hat{P}(x)\hat{P}(y)},$$

where

$$\hat{P}(x, y) = \frac{n(x, y)}{N} \quad \hat{P}(x) = \frac{n(x)}{N},$$

and  $n(x, y)$  denotes the number of records with  $X = x$  and  $Y = y$ .  
Plugging-in these estimates we get:

$$\begin{aligned} I(X; Y) &= \sum_x \sum_y \frac{n(x, y)}{N} \log_2 \frac{n(x, y)/N}{(n(x)/N)(n(y)/N)} \\ &= \sum_x \sum_y \frac{n(x, y)}{N} \log_2 \frac{N \times n(x, y)}{n(x) \times n(y)} \end{aligned}$$

## Estimated Mutual Information

Mutual information between occurrence of the word “bad” and class (negative/positive review):

bad/class	0	1	Total
0	5222	7085	12307
1	2778	915	3693
Total	8000	8000	16000

$$\begin{aligned} I(\text{bad}; \text{class}) &= \frac{5222}{16000} \log_2 \frac{16000 \times 5222}{12307 \times 8000} + \frac{7085}{16000} \log_2 \frac{16000 \times 7085}{12307 \times 8000} \\ &\quad + \frac{2778}{16000} \log_2 \frac{16000 \times 2778}{3693 \times 8000} + \frac{915}{16000} \log_2 \frac{16000 \times 915}{3693 \times 8000} \\ &\approx 0.057 \end{aligned}$$

Fun fact: the estimated mutual information is equal to the deviance of the independence model divided by  $2N$  (if we take the log with base 2 in computing the deviance).

# Movie Reviews: IMDB Review Dataset

- Collection of 50,000 reviews from IMDB, allowing no more than 30 reviews per movie.
- Contains an even number of positive and negative reviews, so random guessing yields 50% accuracy.
- Considers only highly polarized reviews. A negative review has a score  $\leq 4$  out of 10, and a positive review has a score  $\geq 7$  out of 10.
- Neutral reviews are not included in the dataset.

Andrew L. Maas et al., *Learning Word Vectors for Sentiment Analysis*, Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies, pages 142–150, 2011.

Data available at:

<http://ai.stanford.edu/~amaas/data/sentiment/>

# Analysis of Movie Reviews in R

```
# load the tm package
> library(tm)
# Read in the data using UTF-8 encoding
> reviews.neg <- Corpus(DirSource("D:/MovieReviews/train/neg",
                                    encoding="UTF-8"))
> reviews.pos <- Corpus(DirSource("D:/MovieReviews/train/pos",
                                    encoding="UTF-8"))
# Join negative and positive reviews into a single Corpus
> reviews.all <- c(reviews.neg,reviews.pos)
# create label vector (0=negative, 1=positive)
> labels <- c(rep(0,12500),rep(1,12500))
> reviews.all
<<VCorpus>>
Metadata: corpus specific: 0, document level (indexed): 0
Content: documents: 25000
```

# Analysis of Movie Reviews

The first review before pre-processing:

```
> as.character(reviews.all[[1]])  
[1] "Story of a man who has unnatural feelings for a pig.  
Starts out with a opening scene that is a terrific example  
of absurd comedy. A formal orchestra audience is turned into  
an insane, violent mob by the crazy chantings of it's singers.  
Unfortunately it stays absurd the WHOLE time with no  
general narrative eventually making it just too off putting.  
Even those from the era should be turned off.  
The cryptic dialogue would make Shakespeare seem easy to a  
third grader. On a technical level it's better than you might  
think with some good cinematography by future great Vilmos Zsigmond.  
Future stars Sally Kirkland and Frederic Forrest can be seen briefly."
```

# Analysis of Movie Reviews: Pre-Processing

```
# Remove punctuation marks (comma's, etc.)  
> reviews.all <- tm_map(reviews.all,removePunctuation)  
# Make all letters lower case  
> reviews.all <- tm_map(reviews.all,content_transformer(tolower))  
# Remove stopwords  
> reviews.all <- tm_map(reviews.all, removeWords,  
                           stopwords("english"))  
# Remove numbers  
> reviews.all <- tm_map(reviews.all,removeNumbers)  
# Remove excess whitespace  
> reviews.all <- tm_map(reviews.all,stripWhitespace)
```

Not done: stemming, part-of-speech tagging, ...

# Analysis of Movie Reviews

The first review after pre-processing:

```
> as.character(reviews.all[[1]])  
[1] "story man unnatural feelings pig starts opening scene terrific  
example absurd comedy formal orchestra audience turned insane violent  
mob crazy chantings singers unfortunately stays absurd whole time  
general narrative eventually making just putting even era turned  
cryptic dialogue make shakespeare seem easy third grader technical  
level better might think good cinematography future great vilmos  
zsigmond future stars sally kirkland frederic forrest can seen briefly"
```

# Analysis of Movie Reviews

```
# draw training sample (stratified)
# draw 8000 negative reviews at random
> index.neg <- sample(12500,8000)
# draw 8000 positive reviews at random
> index.pos <- 12500+sample(12500,8000)
> index.train <- c(index.neg,index.pos)

# create document-term matrix from training corpus
> train.dtm <- DocumentTermMatrix(reviews.all[index.train])
> dim(train.dtm)
[1] 16000 92564
```

We've got 92,564 features. Perhaps this is a bit too much.

```
# remove terms that occur in less than 5% of the documents
# (so-called sparse terms)

> train.dtm <- removeSparseTerms(train.dtm,0.95)
> dim(train.dtm)
[1] 16000    308
```

# Analysis of Movie Reviews

```
# view a small part of the document-term matrix
> inspect(train.dtm[100:110,80:85])
<<DocumentTermMatrix (documents: 11, terms: 6)>>
Non-/sparse entries: 7/59
Sparsity           : 89%
Maximal term length: 6
Weighting          : term frequency (tf)
```

## Terms

Docs	fact	family	fan	far	father	feel
5663_3.txt	0	0	0	0	0	0
8566_3.txt	0	0	0	0	0	1
10336_4.txt	0	0	0	0	0	0
922_1.txt	0	0	0	1	0	0
8447_3.txt	0	0	0	1	0	0
4062_1.txt	0	0	0	0	0	0
6334_1.txt	0	0	0	0	0	0
333_3.txt	1	0	0	0	0	0
10241_1.txt	0	0	0	0	0	0
831_2.txt	0	0	0	1	0	1
8008_1.txt	0	0	0	0	0	1

# Multinomial naive Bayes in R: Training

```
> train.mnb
function (dtm,labels)
{
call <- match.call()
V <- ncol(dtm)
N <- nrow(dtm)
prior <- table(labels)/N
labelnames <- names(prior)
nclass <- length(prior)
cond.probs <- matrix(nrow=V,ncol=nclass)
dimnames(cond.probs)[[1]] <- dimnames(dtm)[[2]]
dimnames(cond.probs)[[2]] <- labelnames
index <- list(length=nclass)
for(j in 1:nclass){
  index[[j]] <- c(1:N)[labels == labelnames[j]]
}
for(i in 1:V){
  for(j in 1:nclass){
    cond.probs[i,j] <- (sum(dtm[index[[j]],i])+1)/(sum(dtm[index[[j]],])+V)
  }
}
list(call=call,prior=prior,cond.probs=cond.probs)
```

# Multinomial naive Bayes in R: Prediction

```
> predict.mnb
function (model,dtm)
{
  classlabels <- dimnames(model$cond.probs)[[2]]
  logprobs <- dtm %*% log(model$cond.probs)
  N <- nrow(dtm)
  nclass <- ncol(model$cond.probs)
  logprobs <- logprobs+matrix(nrow=N,ncol=nclass,log(model$prior),byrow=T)
  classlabels[max.col(logprobs)]
}
```

# Application of Multinomial naive Bayes to Movie Reviews

```
# Train multinomial naive Bayes model

> reviews.mnb <- train.mnb(as.matrix(train.dtm),labels[index.train])

# create document term matrix for test set
> test.dtm <- DocumentTermMatrix(reviews.all[-index.train],
           list(dictionary=dimnames(train.dtm)[[2]]))
> dim(test.dtm)
[1] 9000 308

> reviews.mnb.pred <- predict.mnb(reviews.mnb,as.matrix(test.dtm))
> table(reviews.mnb.pred,labels[-index.train])

reviews.mnb.pred     0     1
                  0 3459 848
                  1 1041 3652

# compute accuracy on test set: about 79% correct
> (3459+3652)/9000
[1] 0.7901111
```

# Feature Selection with Mutual Information

The top-10 features (terms) according to mutual information are:

term	MI(term, class)
bad	0.044
worst	0.035
great	0.021
awful	0.020
excellent	0.015
terrible	0.014
stupid	0.013
boring	0.012
wonderful	0.012
best	0.011

# Computing Mutual Information

```
# load library "entropy"
> library(entropy)

# compute mutual information of each term with class label
> train.mi <- apply(as.matrix(train.dtm),2,
  function(x,y){mi.plugin(table(x,y)/length(y))},labels[index.train])

# sort the indices from high to low mutual information
> train.mi.order <- order(train.mi,decreasing=T)

# show the five terms with highest mutual information
> train.mi[train.mi.order[1:5]]
      bad      worst     great      awful     excellent
0.04411801 0.03532721 0.02101794 0.01957109 0.01502969
```

# Using the top-50 features

```
# train on the 50 best features
> revs.mnb.top50 <- train.mnb(as.matrix(train.dtm)[,train.mi.order[1:50]],
                                labels[index.train])
# predict on the test set
> revs.mnb.top50.pred <- predict.mnb(revs.mnb.top50,
                                         as.matrix(test.dtm)[,train.mi.order[1:50]])
# show the confusion matrix
> table(revs.mnb.top50.pred,labels[-index.train])

revs.mnb.top50.pred      0      1
                      0 3487  957
                      1 1013 3543

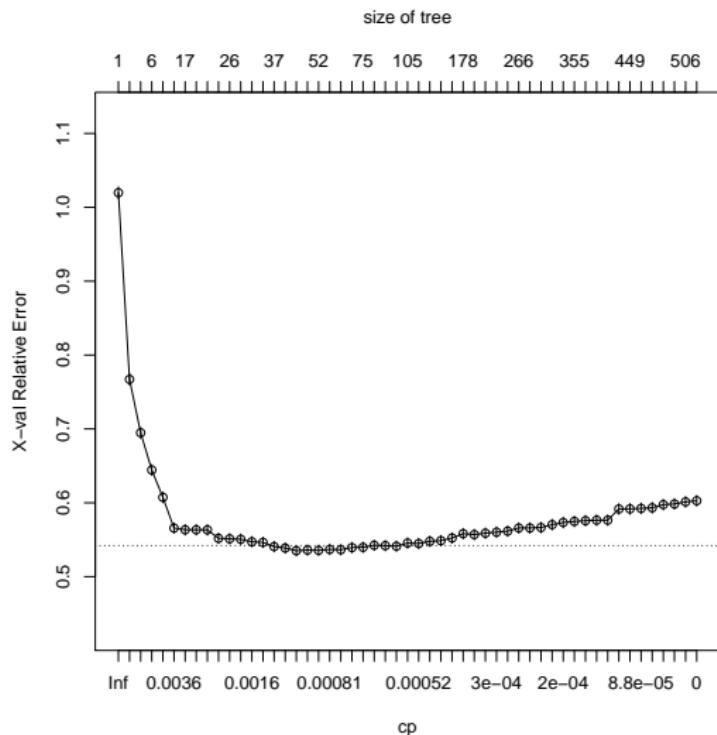
# accuracy is a bit worse compared to using all features
> (3487+3543)/9000
[1] 0.7811111
```

# Classification Trees

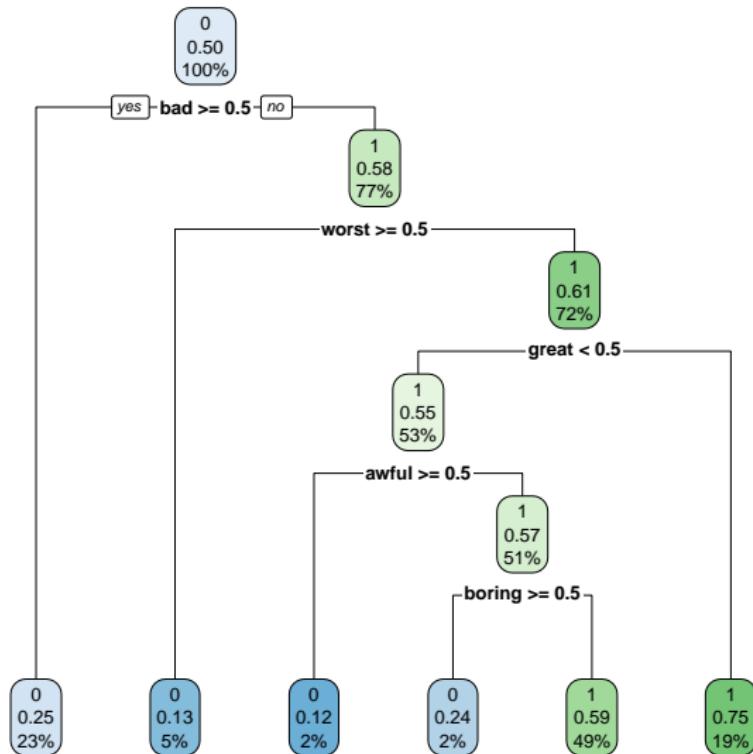
```
# load the required packages
> library(rpart)
> library(rpart.plot)
# grow the tree
> reviews.rpart <- rpart(label~.,
  data=data.frame(as.matrix(train.dtm),label=labels[index.train]),
  cp=0,method="class")
# simple tree for plotting
> reviews.rpart.pruned <- prune(reviews.rpart,cp=1.37e-02)
> rpart.plot(reviews.rpart.pruned)
# tree with lowest cv error
> reviews.rpart.pruned <- prune(reviews.rpart,cp=0.001)
# make predictions on the test set
> reviews.rpart.pred <- predict(reviews.rpart.pruned,
  newdata=data.frame(as.matrix(test.dtm)),type="class")
# show confusion matrix
> table(reviews.rpart.pred,labels[-index.train])

reviews.rpart.pred      0      1
              0 3148 1074
              1 1352 3426
# accuracy is worse than naive Bayes!
> (3148+3426)/9000
[1] 0.7304444
```

# Pruning Sequence



# The Tree



## The Second Assignment: Text Classification

Text Classification for the Detection of Opinion Spam.

- We analyze fake and genuine hotel reviews.
- The genuine reviews have been collected from several popular online review communities.
- The fake reviews have been obtained from Mechanical Turk.
- There are 400 reviews in each of the categories: positive truthful, positive deceptive, negative truthful, negative deceptive.
- We will focus on the negative reviews and try to discriminate between truthful and deceptive reviews.
- Hence, the total number of reviews in our data set is 800.

# The Second Assignment: Text Classification

Analyse the data with:

- ① Naive Bayes (generative linear classifier),
- ② Regularized logistic regression (discriminative linear classifier),
- ③ Classification trees, (flexible classifier) and
- ④ Random forests (flexible classifier).

## The Second Assignment: Text Classification

- This is a data analysis assignment, not a programming assignment.
- You will need to program a little to be able to perform the experiments.
- You only need to hand in a report of your analysis.
- We recommend and support certain R packages (such as `tm`), but you are free to use whatever tools you want.
- For possibly relevant R packages, have a look at:  
<https://cran.r-project.org/web/views/NaturalLanguageProcessing.html>
- The report should describe the analysis you performed in such a way that the reader would be able reproduce it.