

Data Mining 2018

Bayesian Networks (2)

Ad Feelders

Universiteit Utrecht

Learning Bayesian Networks

- ① Parameter learning: structure known/given; we only need to estimate the conditional probabilities from the data.
- ② Structure learning: structure unknown; we need to learn the networks structure as well as the corresponding conditional probabilities from the data.

Smoothing by adding “prior counts”

Recall that the maximum likelihood estimate of $p(x_i | x_{pa(i)})$ is:

$$\hat{p}(x_i | x_{pa(i)}) = \frac{n(x_i, x_{pa(i)})}{n(x_{pa(i)})}$$

But sometimes we have no ($n(x_{pa(i)}) = 0$) or very few observations to estimate these (conditional) probabilities.

Smoothing by adding “prior counts”

Add “prior counts” to “smooth” the estimates.

$$\hat{p}^s(x_i | x_{pa(i)}) = \frac{n(x_{pa(i)})\hat{p}(x_i | x_{pa(i)}) + m(x_{pa(i)})p^0(x_i | x_{pa(i)})}{n(x_{pa(i)}) + m(x_{pa(i)})}$$

where $m(x_{pa(i)})$ is the prior precision, $\hat{p}^s(x_i | x_{pa(i)})$ is the smoothed estimate, and $p^0(x_i | x_{pa(i)})$ is our prior estimate of $p(x_i | x_{pa(i)})$.

Common to take $m(x_{pa(i)})$ to be the same for all parent configurations.

Weighted average of ML estimate and prior estimate.

ML estimate vs. smoothed estimate

For example

$$\hat{p}_{3|1,2}(1|1,2) = \frac{n(x_1 = 1, x_2 = 2, x_3 = 1)}{n(x_1 = 1, x_2 = 2)} = \frac{0}{2} = 0$$

Suppose we set $m = 2$, and $p^0(x_i | x_{pa(i)}) = \hat{p}(x_i)$

Then we get

$$\hat{p}_{3|1,2}^s(1|1,2) = \frac{2 \times 0 + 2 \times 0.4}{2 + 2} = 0.2,$$

since $\hat{p}_3(1) = 0.4$.

Data Set and Likelihood

$$P(X_1, X_2, X_3, X_4) = p_1(X_1)p_2(X_2)p_{3|12}(X_3|X_1, X_2)p_{4|3}(X_4|X_3)$$

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$p_1(1)p_2(1)p_{3|12}(1|1, 1)p_{4|3}(1|1)$$

Data Set and Likelihood

$$P(X_1, X_2, X_3, X_4) = p_1(X_1)p_2(X_2)p_{3|12}(X_3|X_1, X_2)p_{4|3}(X_4|X_3)$$

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$p_1(1)p_2(1)p_{3|12}(1|1, 1)p_{4|3}(1|1)$$
$$p_1(1)p_2(1)p_{3|12}(1|1, 1)p_{4|3}(1|1)$$

Data Set and Likelihood

$$P(X_1, X_2, X_3, X_4) = p_1(X_1)p_2(X_2)p_{3|12}(X_3|X_1, X_2)p_{4|3}(X_4|X_3)$$

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\begin{aligned} & p_1(1)p_2(1)p_{3|12}(1|1, 1)p_{4|3}(1|1) \\ & p_1(1)p_2(1)p_{3|12}(1|1, 1)p_{4|3}(1|1) \\ & p_1(1)p_2(1)p_{3|12}(2|1, 1)p_{4|3}(1|2) \end{aligned}$$

Data Set and Likelihood

$$P(X_1, X_2, X_3, X_4) = p_1(X_1)p_2(X_2)p_{3|12}(X_3|X_1, X_2)p_{4|3}(X_4|X_3)$$

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$p_1(1)p_2(1)p_{3|12}(1|1, 1)p_{4|3}(1|1)$$

$$p_1(1)p_2(1)p_{3|12}(1|1, 1)p_{4|3}(1|1)$$

$$p_1(1)p_2(1)p_{3|12}(2|1, 1)p_{4|3}(1|2)$$

$$p_1(1)p_2(2)p_{3|12}(2|1, 2)p_{4|3}(1|2)$$

Data Set and Likelihood

$$P(X_1, X_2, X_3, X_4) = p_1(X_1)p_2(X_2)p_{3|12}(X_3|X_1, X_2)p_{4|3}(X_4|X_3)$$

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$p_1(1)p_2(1)p_{3|12}(1|1, 1)p_{4|3}(1|1)$
 $p_1(1)p_2(1)p_{3|12}(1|1, 1)p_{4|3}(1|1)$
 $p_1(1)p_2(1)p_{3|12}(2|1, 1)p_{4|3}(1|2)$
 $p_1(1)p_2(2)p_{3|12}(2|1, 2)p_{4|3}(1|2)$
 $p_1(1)p_2(2)p_{3|12}(2|1, 2)p_{4|3}(2|2)$

Data Set and Likelihood

$$P(X_1, X_2, X_3, X_4) = p_1(X_1)p_2(X_2)p_{3|12}(X_3|X_1, X_2)p_{4|3}(X_4|X_3)$$

obs	X_1	X_2	X_3	X_4	
1	1	1	1	1	$p_1(1)p_2(1)p_{3 12}(1 1, 1)p_{4 3}(1 1)$
2	1	1	1	1	$p_1(1)p_2(1)p_{3 12}(1 1, 1)p_{4 3}(1 1)$
3	1	1	2	1	$p_1(1)p_2(1)p_{3 12}(2 1, 1)p_{4 3}(1 2)$
4	1	2	2	1	$p_1(1)p_2(2)p_{3 12}(2 1, 2)p_{4 3}(1 2)$
5	1	2	2	2	$p_1(1)p_2(2)p_{3 12}(2 1, 2)p_{4 3}(2 2)$
6	2	1	1	2	$p_1(2)p_2(1)p_{3 12}(1 2, 1)p_{4 3}(2 1)$
7	2	1	2	3	$p_1(2)p_2(1)p_{3 12}(2 2, 1)p_{4 3}(3 2)$
8	2	1	2	3	$p_1(2)p_2(1)p_{3 12}(2 2, 1)p_{4 3}(3 2)$
9	2	2	2	3	$p_1(2)p_2(2)p_{3 12}(2 2, 2)p_{4 3}(3 2)$
10	2	2	1	3	$p_1(2)p_2(2)p_{3 12}(1 2, 2)p_{4 3}(3 1)$

For all observations

Likelihood function for all observations together:

$$\begin{aligned}L(\mathcal{D}) = & p_1(1)^5(1 - p_1(1))^5 p_2(1)^6(1 - p_2(1))^4 p_{3|1,2}(1|1,1)^2(1 - p_{3|1,2}(1|1,1)) \\& (1 - p_{3|1,2}(1|1,2))^2 p_{3|1,2}(1|2,1)(1 - p_{3|1,2}(1|2,1))^2 p_{3|1,2}(1|2,2) \\& (1 - p_{3|1,2}(1|2,2)) p_{4|3}(1|1)^2 p_{4|3}(2|1)(1 - p_{4|3}(1|1) - p_{4|3}(2|1)) \\& p_{4|3}(1|2)^2 p_{4|3}(2|2)(1 - p_{4|3}(1|2) - p_{4|3}(2|2))^3\end{aligned}$$

Or in log form

$$\begin{aligned}\mathcal{L}(\mathcal{D}) = & 5 \log p_1(1) + 5 \log(1 - p_1(1)) + 6 \log p_2(1) + 4 \log(1 - p_2(1)) \\& + 2 \log p_{3|1,2}(1|1,1) + \log(1 - p_{3|1,2}(1|1,1)) \\& + 2 \log(1 - p_{3|1,2}(1|1,2)) + \log p_{3|1,2}(1|2,1) + 2 \log(1 - p_{3|1,2}(1|2,1)) \\& + \log p_{3|1,2}(1|2,2) + \log(1 - p_{3|1,2}(1|2,2)) \\& + 2 \log p_{4|3}(1|1) + \log p_{4|3}(2|1) + \log(1 - p_{4|3}(1|1) - p_{4|3}(2|1)) \\& + 2 \log p_{4|3}(1|2) + \log p_{4|3}(2|2) + 3 \log(1 - p_{4|3}(1|2) - p_{4|3}(2|2))\end{aligned}$$

Structure Learning

The loglikelihood function for a Bayesian Network is:

$$\mathcal{L} = \sum_{i=1}^k \sum_{x_i, x_{pa(i)}} n(x_i, x_{pa(i)}) \log p(x_i | x_{pa(i)})$$

The maximum likelihood estimate of $p(x_i | x_{pa(i)})$ is:

$$\hat{p}(x_i | x_{pa(i)}) = \frac{n(x_i, x_{pa(i)})}{n(x_{pa(i)})}.$$

The Log-Likelihood Score

If we fill in the maximum likelihood estimates in the log-likelihood function we obtain the value of the log-likelihood function evaluated at its maximum:

$$\mathcal{L} = \sum_{i=1}^k \sum_{x_i, x_{pa(i)}} n(x_i, x_{pa(i)}) \log \frac{n(x_i, x_{pa(i)})}{n(x_{pa(i)})}$$

- This is called the log-likelihood score.
- The higher its value, the better the model fits the data.
- The saturated model (complete graph) always has the highest log-likelihood score.
- To avoid overfitting, we must penalize model complexity.

Scoring Functions for Structure Learning

Scoring functions:

- $AIC(M) = \mathcal{L}^M - \dim(M)$.
- $BIC(M) = \mathcal{L}^M - \frac{\log n}{2} \dim(M)$.

where \mathcal{L}^M is the log-likelihood score of model M and $\dim(M)$ is the number of parameters of M .

BIC gives a higher penalty for model complexity ($n > 7$), so tends to lead to less complex models than AIC.

Note: earlier we defined $AIC(M) = 2(\mathcal{L}^{\text{sat}} - \mathcal{L}^M) + 2\dim(M)$. Dividing by -2 and ignoring the constant \mathcal{L}^{sat} gives the current definition.

Optimization Problem

Given

- Training data.
- Scoring function (BIC or AIC).
- Space of possible models (all DAGs).

find the model that maximizes the score.

- Most model search algorithms do not require an a priori ordering of the variables!
- The number of labeled acyclic directed graphs on k nodes is given by the recurrence

$$a_k = \sum_{j=1}^k (-1)^{j-1} \binom{k}{j} 2^{j(k-j)} a_{k-j}$$

For example, $a_6 = 3,781,503$.

Heuristic Search

- Define which models are neighbors of a given model (typically: addition, removal, or reversal of an arc).
- Traverse search space looking for high-scoring models, e.g. by greedy hill-climbing.

Score Decomposes

The loglikelihood score

$$\mathcal{L} = \sum_{i=1}^k \sum_{x_i, x_{pa(i)}} n(x_i, x_{pa(i)}) \log \frac{n(x_i, x_{pa(i)})}{n(x_{pa(i)})}$$

must be computed many times for different models in structure learning.

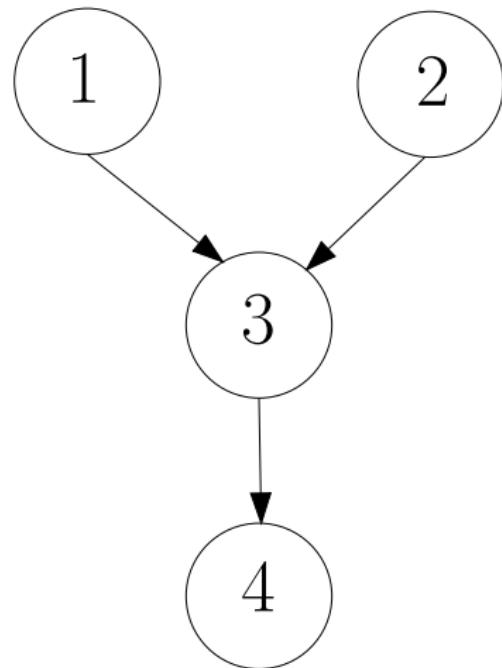
Luckily, it is a sum of terms, where each term contains the variables $\{i\} \cup pa(i)$.

Hence, when making a change to the model, we only have to recompute the score for those variables for which the parent set has changed!

Example Data Set

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

Score this model



Relevant Data For Scoring Node 1

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{Score node 1} = 5 \log \frac{5}{10} + 5 \log \frac{5}{10}$$

Relevant Data For Scoring Node 2

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{Score node 2} = 6 \log \frac{6}{10} + 4 \log \frac{4}{10}$$

Relevant Data For Scoring Node 3

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{Score node 3} = 2 \log \frac{2}{3} + \log \frac{1}{3}$$

Relevant Data For Scoring Node 3

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{Score node 3} = 2 \log \frac{2}{3} + \log \frac{1}{3} + 2 \log 1$$

Relevant Data For Scoring Node 3

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{Score node 3} = 2 \log \frac{2}{3} + \log \frac{1}{3} + 2 \log 1 + \log \frac{1}{3} + 2 \log \frac{2}{3}$$

Relevant Data For Estimating Scoring Node 3

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{Score node 3} = 2 \log \frac{2}{3} + \log \frac{1}{3} + 2 \log 1 + \log \frac{1}{3} + 2 \log \frac{2}{3} + \log \frac{1}{2} + \log \frac{1}{2}$$

Relevant Data For Scoring Node 4

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{Score node 4} = 2 \log \frac{2}{4} + \log \frac{1}{4} + \log \frac{1}{4}$$

Relevant Data For Scoring Node 4

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{Score node 4} = 2 \log \frac{2}{4} + \log \frac{1}{4} + \log \frac{1}{4} + 2 \log \frac{2}{6} + \log \frac{1}{6} + 3 \log \frac{3}{6}$$

Total Score

Summing the loglikelihood score over all nodes, we get:

$$\mathcal{L} = 5 \log \frac{5}{10} + 5 \log \frac{5}{10} \quad (\text{node 1})$$

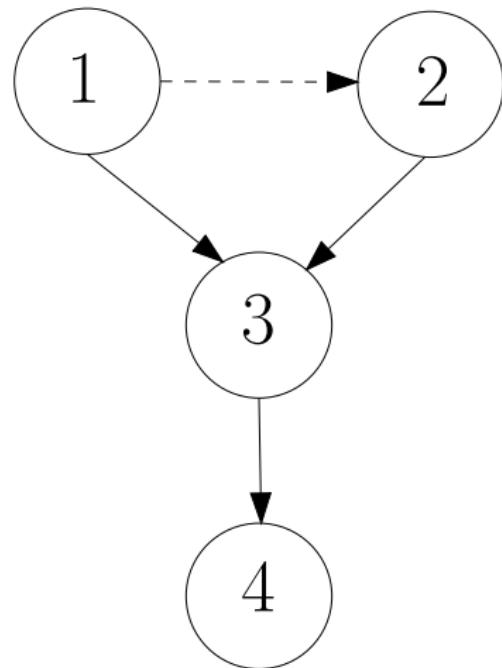
$$+ 6 \log \frac{6}{10} + 4 \log \frac{4}{10} \quad (\text{node 2})$$

$$+ 2 \log \frac{2}{3} + \log \frac{1}{3} + 2 \log 1 + \log \frac{1}{3} + 2 \log \frac{2}{3} + \log \frac{1}{2} + \log \frac{1}{2} \quad (\text{node 3})$$

$$+ 2 \log \frac{2}{4} + \log \frac{1}{4} + \log \frac{1}{4} + 2 \log \frac{2}{6} + \log \frac{1}{6} + 3 \log \frac{3}{6} \quad (\text{node 4})$$

$$\approx -29.09$$

Add an edge from X_1 to X_2



Score is Decomposable

$$\begin{aligned}\mathcal{L} &= 5 \log \frac{5}{10} + 5 \log \frac{5}{10} && (\text{node 1}) \\ &+ \boxed{6 \log \frac{6}{10} + 4 \log \frac{4}{10}} && (\text{node 2}) \\ &+ 2 \log \frac{2}{3} + \log \frac{1}{3} + 2 \log 1 + \log \frac{1}{3} + 2 \log \frac{2}{3} + \log \frac{1}{2} + \log \frac{1}{2} && (\text{node 3}) \\ &+ 2 \log \frac{2}{4} + \log \frac{1}{4} + \log \frac{1}{4} + 2 \log \frac{2}{6} + \log \frac{1}{6} + 3 \log \frac{3}{6} && (\text{node 4}) \\ &\approx -29.09\end{aligned}$$

- When we add an edge from X_1 to X_2 , only the parent set of node 2 changes.
- Therefore, only the score of node 2 (the boxed part) has to be recomputed.

Relevant Data For Re-scoring Node 2

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{New score node 2} = 3 \log \frac{3}{5} + 2 \log \frac{2}{5}$$

Relevant Data For Re-scoring Node 2

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{New score node 2} = 3 \log \frac{3}{5} + 2 \log \frac{2}{5} + 3 \log \frac{3}{5} + 2 \log \frac{2}{5}$$

Score Decomposes

$$\mathcal{L} = 5 \log \frac{5}{10} + 5 \log \frac{5}{10} \quad (\text{node 1})$$

$$+ \boxed{3 \log \frac{3}{5} + 2 \log \frac{2}{5} + 3 \log \frac{3}{5} + 2 \log \frac{2}{5}} \quad (\text{node 2})$$

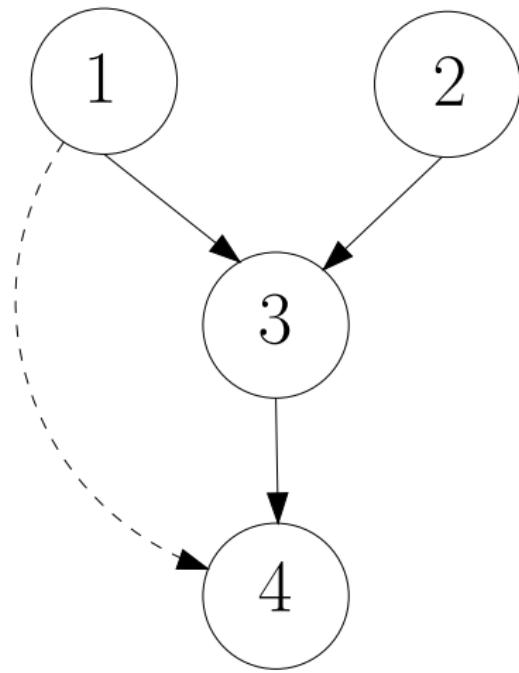
$$+ 2 \log \frac{2}{3} + \log \frac{1}{3} + 2 \log 1 + \log \frac{1}{3} + 2 \log \frac{2}{3} + \log \frac{1}{2} + \log \frac{1}{2} \quad (\text{node 3})$$

$$+ 2 \log \frac{2}{4} + \log \frac{1}{4} + \log \frac{1}{4} + 2 \log \frac{2}{6} + \log \frac{1}{6} + 3 \log \frac{3}{6} \quad (\text{node 4})$$

$$\approx -29.09$$

The boxed part is the new contribution of node 2 to the score.

Add an edge from X_1 to X_4



Score Decomposes

$$\begin{aligned}\mathcal{L} &= 5 \log \frac{5}{10} + 5 \log \frac{5}{10} && (\text{node 1}) \\ &+ 6 \log \frac{6}{10} + 4 \log \frac{4}{10} && (\text{node 2}) \\ &+ 2 \log \frac{2}{3} + \log \frac{1}{3} + 2 \log 1 + \log \frac{1}{3} + 2 \log \frac{2}{3} + \log \frac{1}{2} + \log \frac{1}{2} && (\text{node 3}) \\ &+ \boxed{2 \log \frac{2}{4} + \log \frac{1}{4} + \log \frac{1}{4} + 2 \log \frac{2}{6} + \log \frac{1}{6} + 3 \log \frac{3}{6}} && (\text{node 4}) \\ &\approx -29.09\end{aligned}$$

- When we add an edge from X_1 to X_4 , only the parent set of node 4 changes.
- Therefore, only the score of node 4 (the boxed part) has to be recomputed.

Relevant Data For Re-scoring Node 4

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

New score node 4 = $2 \log 1$

Relevant Data For Re-scoring Node 4

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{New score node 4} = 2 \log 1 + 2 \log \frac{2}{3} + \log \frac{1}{3}$$

Relevant Data For Re-scoring Node 4

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{New score node 4} = 2 \log 1 + 2 \log \frac{2}{3} + \log \frac{1}{3} + \log \frac{1}{2} + \log \frac{1}{2}$$

Relevant Data For Re-scoring Node 4

obs	X_1	X_2	X_3	X_4
1	1	1	1	1
2	1	1	1	1
3	1	1	2	1
4	1	2	2	1
5	1	2	2	2
6	2	1	1	2
7	2	1	2	3
8	2	1	2	3
9	2	2	2	3
10	2	2	1	3

$$\text{New score node 4} = 2 \log 1 + 2 \log \frac{2}{3} + \log \frac{1}{3} + \log \frac{1}{2} + \log \frac{1}{2} + 3 \log 1$$

Score Decomposes

$$\mathcal{L} = 5 \log \frac{5}{10} + 5 \log \frac{5}{10} \quad (\text{node 1})$$

$$+ 6 \log \frac{6}{10} + 4 \log \frac{4}{10} \quad (\text{node 2})$$

$$+ 2 \log \frac{2}{3} + \log \frac{1}{3} + 2 \log 1 + \log \frac{1}{3} + 2 \log \frac{2}{3} + \log \frac{1}{2} + \log \frac{1}{2} \quad (\text{node 3})$$

$$+ \boxed{2 \log 1 + 2 \log \frac{2}{3} + \log \frac{1}{3} + \log \frac{1}{2} + \log \frac{1}{2} + 3 \log 1} \quad (\text{node 4})$$

$$\approx -22.16$$

The boxed part is the new contribution of node 4 to the score.

Counting Parameters

The number of parameters of a Bayesian network is:

$$\sum_{i=1}^k \left\{ (d_i - 1) \prod_{j \in pa(i)} d_j \right\}$$

where k is the number of variables in the network, and d_i is the number of possible values of X_i .

$\prod_{j \in pa(i)} d_j$ is the number of parent configuration for X_i .

If X_i has no parents, the number of parent configurations should be taken to be 1, so X_i contributes $d_i - 1$ parameters in that case.

A Simple Structure Learning Algorithm

Algorithm 1 BN Structure Learning

```
1:  $G \leftarrow$  initial graph
2:  $\max \leftarrow \text{score}(G)$ 
3: repeat
4:    $\text{nb} \leftarrow \text{neighbours}(G)$ 
5:   for all  $G' \in \text{nb}$  do
6:     if  $\text{score}(G') > \max$  then
7:        $\max \leftarrow \text{score}(G')$ 
8:        $G \leftarrow G'$ 
9:     end if
10:   end for
11: until no change to  $G$ 
12: return  $G$ 
```

Complexity: Naive

- A DAG with k nodes has $k(k - 1)$ possible directed edges.

- edge present: delete or reverse
- edge absent: add

So there are $O(k^2)$ neighbours that have to be scored.

- There are k components in the score and for each component we have to compute the sufficient statistics which requires traversing the training data (n rows).

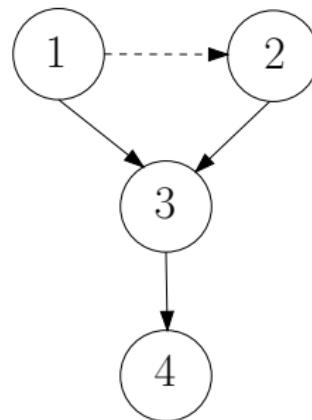
So in total, scoring a single neighbour takes $O(kn)$ time.

- The total complexity is $O(k^3n)$ per search step.

Complexity: Exploiting Decomposability

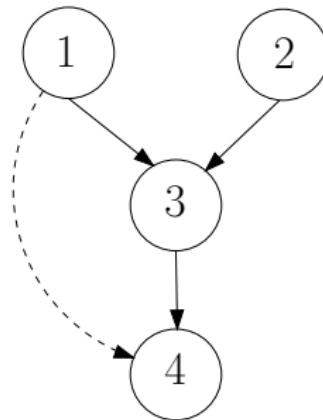
- To score a neighbour we actually only have to score a single node (add, delete), or two nodes (reversal), since we only have to recompute the score for the nodes whose parent set changed by the local operation performed. So the complexity of scoring a neighbour is only $O(n)$ instead of $O(kn)$.
- If we compute the *change* in score due to a local operation (add, delete, reverse), then we can reuse the “delta scores” computed in previous iterations. We only need to recompute the delta scores for nodes whose parent set changed in the previous iteration. So we only have to compute $O(k)$ changes in score.
- Hence the total complexity per search step is $O(kn)$.

Add an edge from X_1 to X_2



$$\begin{aligned}\Delta \text{Score}(\text{add}(X_1 \rightarrow X_2)) &= \left(3 \log \frac{3}{5} + 2 \log \frac{2}{5} + 3 \log \frac{3}{5} + 2 \log \frac{2}{5} \right) \\ &\quad - \left(6 \log \frac{6}{10} + 4 \log \frac{4}{10} \right) = 0\end{aligned}$$

Add an edge from X_1 to X_4

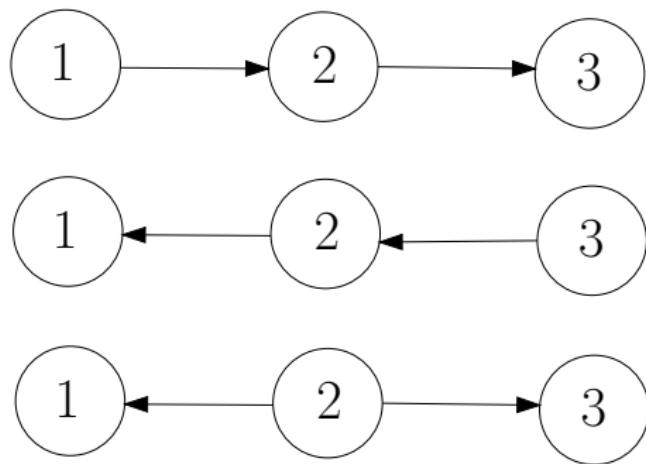


$$\begin{aligned}\Delta \text{Score}(\text{add}(X_1 \rightarrow X_4)) &= \left(2 \log 1 + 2 \log \frac{2}{3} + \log \frac{1}{3} + \log \frac{1}{2} + \log \frac{1}{2} + 3 \log 1 \right) \\ &\quad - \left(2 \log \frac{2}{4} + \log \frac{1}{4} + \log \frac{1}{4} + 2 \log \frac{2}{6} + \log \frac{1}{6} + 3 \log \frac{3}{6} \right) \\ &\approx 6.93\end{aligned}$$

Add an edge from X_1 to X_4

- Suppose we decide to add the arrow $X_1 \rightarrow X_4$.
- In the next iteration only the Δ scores of operations that change the parent set of X_4 have to be recomputed.
- For example, $\Delta\text{Score}(\text{add}(X_1 \rightarrow X_2))$ doesn't have to be recomputed because it is the same as in the previous iteration.

Interpretation: warning!



These models can not be distinguished from data alone.
They represent the same independencies!

AIC and *BIC* give equivalent networks the same score.

Markov Equivalence and Essential Graph

Two DAGs are Markov equivalent if and only if

- ① they have the same skeleton (same undirected graph when you drop the directions of all edges), and
- ② they have the same immoralities (v-structures).

Essential Graph:

For a given DAG, an edge becomes bi-directional in the essential graph if there is an equivalent DAG in which the direction of the edge is reversed.

Example Analysis

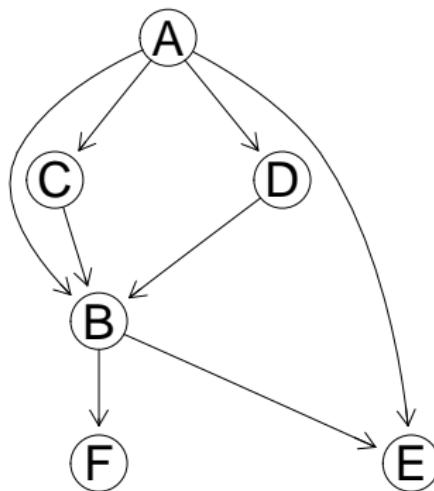
We analyze a data set concerning risk factors for coronary heart disease. For a sample of 1841 car-workers, the following information was recorded

Variable	Description
A	Does the person smoke?
B	Is the person's work strenuous mentally?
C	Is the person's work strenuous physically?
D	Systolic blood pressure < 140mm?
E	Ratio of beta to alfa lipoproteins < 3?
F	Is there a family history of coronary heart disease?

Example Analysis

For learning Bayesian networks, we use the *bnlearn* package in R.
Hill-climbing with the BIC score function:

```
> coronary.hc <- hc(coronary)
> plot(coronary.hc)
```



The Search Process

```
> coronary.hc <- hc(coronary, debug=T)
-----
* starting from the following network:
model:
[A] [B] [C] [D] [E] [F]

* current score: -7061.714
* caching score delta for arc A -> B (17.531166).
* caching score delta for arc A -> C (9.981480).
* caching score delta for arc A -> D (1.757126).
* caching score delta for arc A -> E (4.941129).
* caching score delta for arc A -> F (-3.224701).
* caching score delta for arc B -> C (264.272873).
* caching score delta for arc B -> D (2.313656).
* caching score delta for arc B -> E (21.030213).
* caching score delta for arc B -> F (2.303571).
* caching score delta for arc C -> D (-3.711314).
* caching score delta for arc C -> E (4.577177).
* caching score delta for arc C -> F (-3.673929).
* caching score delta for arc D -> E (2.645583).
* caching score delta for arc D -> F (-3.197133).
* caching score delta for arc E -> F (-2.257169).
```

The Search Process

- The initial model (the mutual independence model [A] [B] [C] [D] [E] [F]) has a BIC score of -7061.714.
- The output gives the *change* in score between the current model and its neighbors.
- Why is the score of only 15 of the 30 neighbors computed?
(e.g. A → B, but not B → A)?

The Search Process

- The initial model (the mutual independence model [A] [B] [C] [D] [E] [F]) has a BIC score of -7061.714.
- The output gives the *change* in score between the current model and its neighbors.
- Why is the score of only 15 of the 30 neighbors computed? (e.g. A → B, but not B → A)?
- A → B and B → A are Markov equivalent, and therefore have the same score.
- Adding B → C results in the largest increase in score so we move to that neighbor.

The Search Process

- * best operation was: adding $B \rightarrow C$.
- * current network is :

model:

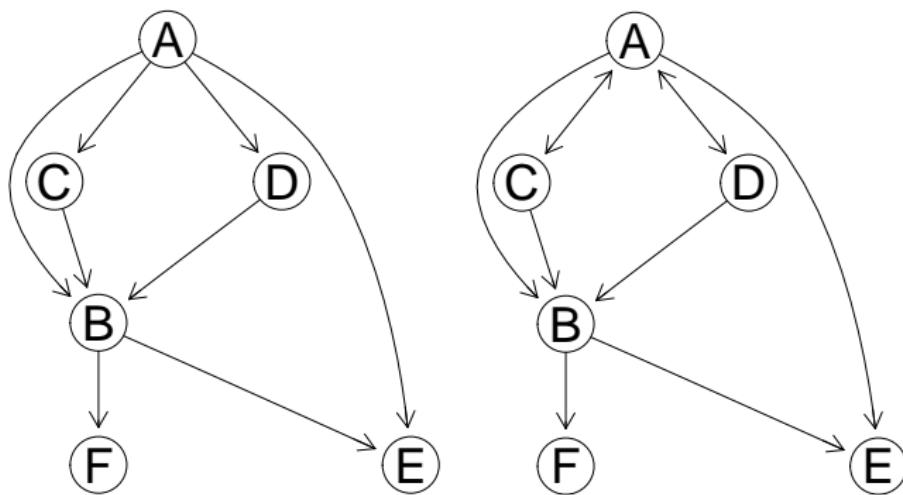
[A] [B] [D] [E] [F] [C|B]

- * current score: -6797.441
- * caching score delta for arc $A \rightarrow C$ (9.975823) .
- * caching score delta for arc $B \rightarrow C$ (-264.272873) .
- * caching score delta for arc $D \rightarrow C$ (-1.472731) .
- * caching score delta for arc $E \rightarrow C$ (-6.587044) .
- * caching score delta for arc $F \rightarrow C$ (-6.059896) .

The Search Process

- We don't have to recompute the change in score caused by, for example, adding $A \rightarrow B$, because the parent set of B is the same as in the previous iteration.
- Therefore, adding $A \rightarrow B$ now will cause the same score change as in the previous iteration.
- Only the parent set of C has changed, so we just have to recompute the change in score caused by adding arcs $X \rightarrow C$.
- Adding $B \rightarrow E$ results in the largest increase in score so we move to that neighbor.
- The current model becomes: $[A] [B] [D] [F] [C|B] [E|B]$.

Final Model and its Essential Graph



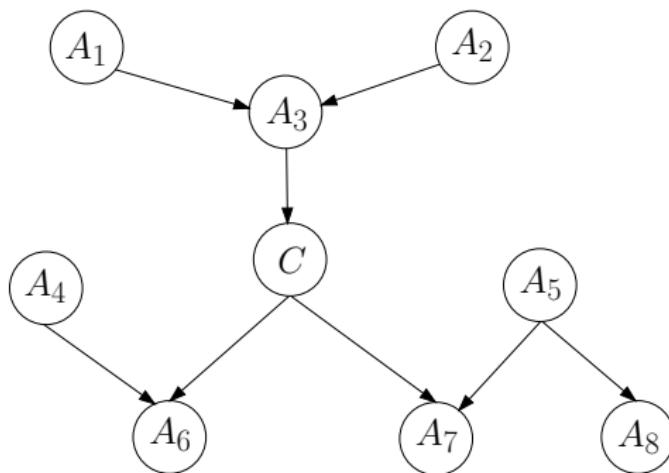
Final Model

Essential Graph

Example of Model Use

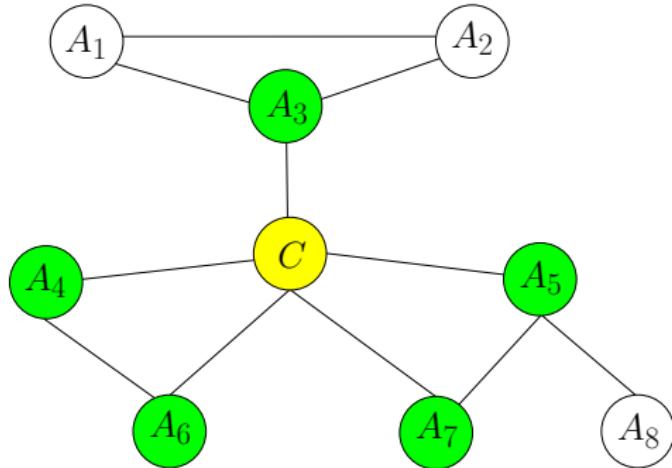
```
# estimate parameters for selected model structure
> coronary.hc.fit <- bn.fit(coronary.hc,coronary.dat,"mle")
# predict B from remaining variables
> coronary.hc.pred <- predict(coronary.hc.fit,node="B",
                                 data=coronary.dat)
# make confusion matrix
> table(coronary.dat$B,coronary.hc.pred)
coronary.hc.pred
      no yes
no   944 186
yes  208 503
> (944+503)/1841
[1] 0.7859859
> (944+186)/1841
[1] 0.6137968
```

Bayesian Networks as Classifiers



Markov Blanket: Parents, Children and Parents of Children.

Markov Blanket of C : Moral Graph



Markov Blanket: Parents, Children and Parents of Children.

Local Markov property: $C \perp\!\!\!\perp \text{rest} \mid \text{boundary}(C)$

Right Heart Catheterization Data: Variable Description

- ① cat1: primary disease category (9 different values)
- ② death: did the patient die within 180 days after admission to ICU?
- ③ swang1: was right heart catheterization (Swan-Ganz catheter) performed within first 24 hours?
- ④ gender: male/female
- ⑤ race: black/white/other
- ⑥ ninsclas: type of medical insurance of patient (six different values)
- ⑦ income: income of patient, divided into 4 categories
- ⑧ ca: cancer status (yes/no/metastatic)
- ⑨ age: age of patient divided into 5 categories
- ⑩ meanbp1: mean blood pressure of patient divided into 2 categories

Right Heart Catheterization Data: Descriptive Statistics

```
> summary(rhc.dat)
```

	cat1	death	swang1	gender	race
ARF	:2490	No :2013	No RHC:3551	Female:2543	black: 920
MOSF w/Sepsis	:1227	Yes:3722	RHC :2184	Male :3192	other: 355
COPD	: 457				white:4460
CHF	: 456				
Coma	: 436				
MOSF w/Malignancy	: 399				
(Other)	: 270				

	ninsclas	income	ca	age
Medicaid	: 647	\$11-\$25k :1165	Metastatic: 384	(50,60] : 917
Medicare	:1458	\$25-\$50k : 893	No :4379	(60,70] :1390
Medicare & Medicaid	: 374	> \$50k : 451	Yes : 972	(70,80] :1337
No insurance	: 322	Under \$11k:3226		(80,102]: 667
Private	:1698			[18,50] :1424
Private & Medicare	:1236			

meanbp1
(85,259]:1975
[0,85] :3760

Learning the Graph Structure

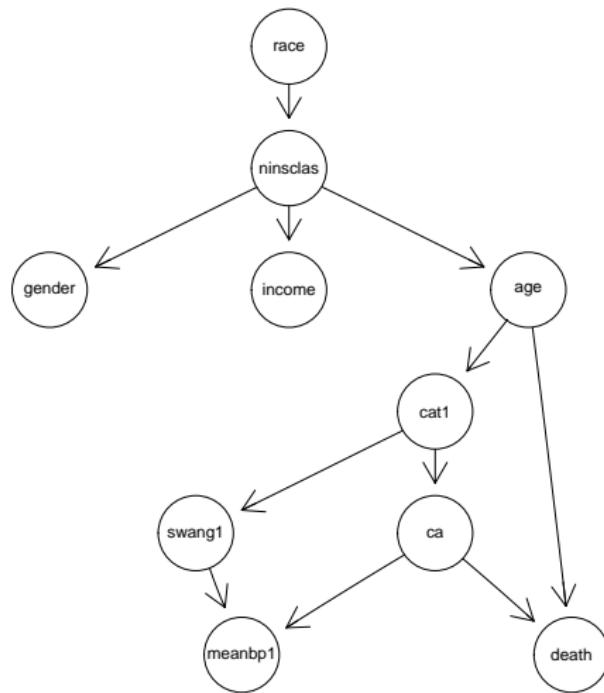
```
# load bayesian network library
> library(bnlearn)

# load library for graph vizualization
> library(Rgraphviz)

# use hill climbing with BIC scoring
# starting from empty graph
> rhc.bn <- hc(rhc.dat)

# plot the model structure
> plot(as(amat(rhc.bn), "graphNEL"))
```

Graph Structure for Right Heart Catheterization Data



Performing Inference

```
# estimate the network parameters
> rhc.bn.fit <- bn.fit(rhc.bn,data=rhc.dat)

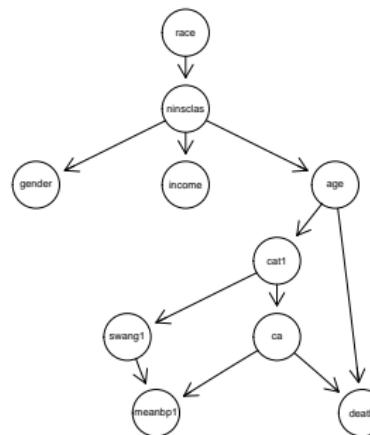
# perform sampling based inference
# probability of death for metastatic cancer and
# mean blood pressure > 85
> cpquery(rhc.bn.fit,event=death=="Yes",evidence=
  ca=="Metastatic" & meanbp1=="(85,259]",n=100000)
[1] 0.9033019

# probability of death for no cancer and
# mean blood pressure > 85
> cpquery(rhc.bn.fit,event=death=="Yes",evidence=
  ca=="No" & meanbp1=="(85,259]",n=100000)
[1] 0.6020206
```

Combining Data and Prior Knowledge

An expert studies the graph and argues that the edge from `swang1` to `meanbp1` is in the wrong direction, since the blood pressure influences the decision to apply right heart catheterization, not the other way around.

Can we turn the edge around without changing the “meaning” of the network, i.e. without changing the conditional independencies expressed by the graph?



Combining Data and Prior Knowledge

Common sense suggest that the variables can be divided into a number of ordered blocks, where arrows are not allowed to point from a variable in a higher block to a variable in a lower block.

As an example, consider the following block structure:

- ① race, gender
- ② age, income
- ③ ninsclass
- ④ cat1, ca, meanbp1
- ⑤ swang1
- ⑥ death

We can use the `blacklist` parameter to avoid edges pointing from higher blocks to lower blocks.

Learning with a Blacklist

```
# learn structure with blacklist
> rhc.bn.ord <- hc(rhc.dat,blacklist=blackL)

# has the score become much worse?
> score(rhc.bn.ord,rhc.dat)
[1] -54059.03
> score(rhc.bn,rhc.dat)
[1] -53749.15

# has inferences changed much?
> rhc.bn.ord.fit <- bn.fit(rhc.bn.ord,data=rhc.dat)

> cpquery(rhc.bn.ord.fit,event=death=="Yes",evidence=
  ca=="Metastatic" & meanbp1=="(85,259]",n=100000)
[1] 0.9039467
> cpquery(rhc.bn.ord.fit,event=death=="Yes",evidence=
  ca=="No" & meanbp1=="(85,259]",n=100000)
[1] 0.610249
```

The Blacklist

The blacklist simply enumerates all the forbidden edges:

```
> blackL
```

	X1	X2
1	cat1	gender
2	cat1	race
3	cat1	ninsclas
4	cat1	income
5	cat1	age
6	death	cat1
7	death	swang1

etc.

Graph Structure Learned with Blacklist

