


Data Science & Society

Lecture 04: *Spark - Architecture & Transformations*

INFOMDSS 2018 :: Dr. Marco Spruit






Agenda

- › Book review Q/A
- › Spark
 - High-level Architecture
 - The concept of Resilient Distributed Datasets (RDDs)
 - Narrow Transformations in Spark
 - ~~Wide Transformations in Spark~~

Architecture of Spark

A high-level outlook onto its components




From lecture 03


Recap: Shortcomings of MapReduce

1. Force your pipeline into Map and Reduce steps
 - Other workflows? i.e. join, filter, map-reduce-map
2. Read from disk for each MapReduce job
 - Iterative algorithms? i.e. machine learning
3. Only native JAVA programming interface
 - Other languages? Interactivity?

How Java's Floating-Point Hurts Everyone Everywhere



From lecture 03




Recap: Solutions by Spark

1. Force your pipeline into Map and Reduce steps
 - Other workflows? i.e. join, filter, map-reduce-map
 - 20 highly efficient distributed operations, any combination of them
2. Read from disk for each MapReduce job
 - Iterative algorithms? i.e. machine learning
 - in-memory caching of data, specified by the user
3. Only native JAVA programming interface
 - Other languages? Interactivity?
 - Native Python, Scala (, R) interface. Interactive shells

Recap: Original MapReduce Execution Framework


- › Software framework which
 - Schedules, monitors, and manages tasks
- › Works for Applications that fit MapReduce paradigm



One `map()` executor
per worker node




Spark Architecture: One Worker node







Multiple Worker nodes






Multiple Worker nodes, managed







Driver Program




From lecture 02

Hadoop's YARN






On a local VM






On Azure, or Amazon Elastic MapReduce (EMR)



Elastic Compute Cloud (EC2) nodes

Resilient Distributed Datasets

RDDs in Apache Spark





What's in a word? Resilient Distributed Dataset

Dataset

Data storage created from: HDFS, S3, HBase, JSON, text,
Local hierarchy of folders

Or created transforming another RDD

i.e. You cannot change just a chunk of them
RDDs are immutable:



What's in a word? Resilient Distributed Dataset

Distributed

Distributed across the cluster of machines

Divided in partitions, atomic chunks of data


i.e. Not just fixed line by line processing
Spark works with partitions

What's in a word? Resilient Distributed Dataset


Resilient

Recover from errors, e.g. node failure, slow processes

Track history of each partition, re-run



(Narrow) Transformations in Apache Spark





Transformations

- › RDD are immutable
- › Never modify RDD in place
- › Transform RDD to another RDD
- › Lazy

Immutable object

msdn.microsoft.com

In object-oriented and functional programming, an immutable object is an object whose state cannot be modified after it is created. This is in contrast to a mutable object, which can be modified after it is created. [Meer op Wikipedia](#)



Some available transformations

- › `map(func)`
 - apply function to each element of RDD
- › `flatMap(func)`
 - map then flatten output
- › `filter(func)`
 - keep only elements where func is true
- › `sample(withReplacement, fraction, seed)`
 - get a random data fraction
- › `coalesce(numPartitions)`
 - merge partitions to reduce them to numPartitions



Example transformation: map(func)

1. Create RDD

- from local filesystem:

```
text_RDD = sc.textFile("file:///home/cloudera/testfile1")
```

```
text_RDD.collect()
```

```
Out[n]: [u'A long time ago in a galaxy far far away']
```

2. Define a mapper function


```
def lower(line):  
    return line.lower()
```

3. Invoke the transformation: map

```
lower_text_RDD = text_RDD.map(lower)
```



Example transformation: map(func)



- › Output size: unchanged




Example transformation: flatMap(func)

1. Define a mapper function

```
def split_words(line):  
    return line.split()
```

2. Invoke the transformation: flatMap

```
words_RDD = text_RDD.flatMap(split_words)
```





Comparison: testfile1 after split_words with...

› Map

```
[[u'A',
  u'long',
  u'time',
  u'ago',
  u'in',
  u'a',
  u'galaxy',
  u'far',
  u'far',
  u'away']]
```

› FlatMap

```
[u'A',
  u'long',
  u'time',
  u'ago',
  u'in',
  u'a',
  u'galaxy',
  u'far',
  u'far',
  u'away']]
```




Example transformation: filter(func)

1. Define a function to filter out all words which start with an "a"

```
def starts_with_a(word):  
    return word.lower().startswith("a")
```

```
words_RDD.filter(starts_with_a).collect()
```

```
Out[ ]:[u'A', u'ago', u'a', u'away']
```



filter() can result in unevenly distributed partitions...




Example transformation: coalesce(func)

1. Create a test array with numbers 1-10 in 4 partitions


```
sc.parallelize(range(10), 4).glom().collect()  
Out[ ]: [[0,1],[2,3],[4,5],[6,7,8,9]]
```


2. Coalesce returns a new RDD that is reduced into 2 partitions

```
sc.parallelize(range(10), 4).coalesce(2).glom().collect()  
Out[ ]: [[0,1,2,3],[4,5,6,7,8,9]]
```




Spark Demo





Initial Wordcount code in Python (for PySpark)

```
Bash      ▾ | ⌂ ? ⚙ ⌂+ {}  
  
def split_words(line):  
    return line.split()  
def create_pair(word):  
    return (word,1)  
def sum_counts(a,b):  
    return a + b  
  
from pyspark.sql import SparkSession  
spark = SparkSession.builder.appName("appName").getOrCreate()  
sc = spark.sparkContext  
  
text_RDD = sc.textFile("file:///home/marco/testfile")  
pairs_RDD = text_RDD.flatMap(split_words).map(create_pair)  
wordcounts_RDD = pairs_RDD.reduceByKey(sum_counts)  
  
wordcounts_RDD.saveAsTextFile("file:///home/marco/output");
```



```
marco@ads:~$ more testfile
A long long long time ago,
in a galaxy far far away...
marco@ads:~$
```


Input > Output

```
marco@ads:~$ ll output/
total 28
drwxrwxr-x  2 marco marco 4096 Sep 19 21:23 .
drwxr-xr-x 12 marco marco 4096 Sep 19 21:23 ../
-rw-r--r--  1 marco marco   60 Sep 19 21:23 part-00000
-rw-r--r--  1 marco marco   12 Sep 19 21:23 .part-00000.crc
-rw-r--r--  1 marco marco   53 Sep 19 21:23 part-00001
-rw-r--r--  1 marco marco   12 Sep 19 21:23 .part-00001.crc
-rw-r--r--  1 marco marco     0 Sep 19 21:23 _SUCCESS
-rw-r--r--  1 marco marco     8 Sep 19 21:23 ._SUCCESS.crc
marco@ads:~$ more output/part-00000
(u'A', 1)
(u'far', 2)
(u'galaxy', 1)
(u'a', 1)
(u'long', 3)
marco@ads:~$ more output/part-00001
(u'ago,', 1)
(u'away...', 1)
(u'in', 1)
(u'time', 1)
marco@ads:~$
```



Updated Wordcount code in Python (for PySpark)

```
Bash    ▾ | ⚡ ? 🛡 { }  
  
def split_words(line):  
    return line.split()  
def create_pair(word):  
    return (word,1)  
def sum_counts(a,b):  
    return a + b  
  
from pyspark.sql import SparkSession  
spark = SparkSession.builder.appName("appName").getOrCreate()  
sc = spark.sparkContext  
  
text_RDD = sc.textFile("file:///home/marco/testfile")  
pairs_RDD = text_RDD.flatMap(split_words).map(create_pair)  
wordcounts_RDD = pairs_RDD.reduceByKey(sum_counts, numPartitions=1)  
  
wordcounts_RDD.saveAsTextFile("file:///home/marco/output");
```




Commands used

- › vi wordcount_mapreduce.py
- › chmod +x wordcount_mapreduce.py
- › spark-submit wordcount_mapreduce.py
- › more output/part-00000
- › rm -r ./output

Optionally: interactive tutorial at:
<http://learnpython.org>

PySpark reference:
<https://spark.apache.org/docs/latest/api/python/pyspark.html>



Final Wordcount output

```
marco@ads:~$ ll output/
total 20
drwxrwxr-x  2 marco marco 4096 Sep 19 21:28 ./
drwxr-xr-x 12 marco marco 4096 Sep 19 21:28 ../
-rw-r--r--  1 marco marco   13 Sep 19 21:28 part-00000
-rw-r--r--  1 marco marco    12 Sep 19 21:28 .part-00000.crc
-rw-r--r--  1 marco marco     0 Sep 19 21:28 _SUCCESS
-rw-r--r--  1 marco marco     8 Sep 19 21:28 ._SUCCESS.crc
marco@ads:~$ more output/part-00000
(u'A', 1)
(u'ago,', 1)
(u'far', 2)
(u'in', 1)
(u'long', 3)
(u'away...', 1)
(u'a', 1)
(u'time', 1)
(u'galaxy', 1)
marco@ads:~$
```