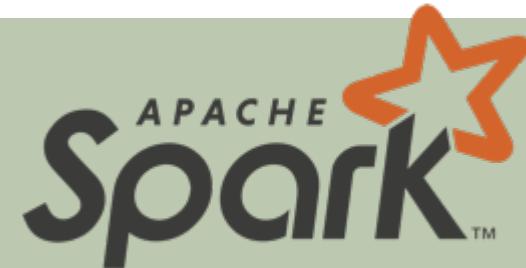
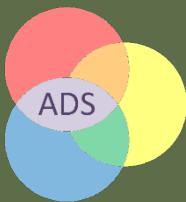


Data Science & Society

Lecture 07: SQL & Spark

INFOMDSS 2018 :: Dr. Marco Spruit





Agenda

- › Book review grades
 - › Tweaked Course Schedule
 - <http://bit.ly/infomdss-schedule>
 - › This Thursday 14:00-14:30: Office hour Midterm exam review
 - › Extra assignment Big Spatial Data
-
- › SQL in context and examples
 - NoSQL, RDD, ...
 - › SQL in Spark
 - CH 10 intro



Course Schedule 2018

<http://bit.ly/infomdss-schedule>

infomdss2018	Date	Monday C3	Tuesday C4 Workshops	Tuesday C5 Lectures	Thursday C7 lectures	Weekly assignments	Weekly readings	Dev/Test Lab
week 1 (37)	2018-09-03		N/A	N/A	Regular lecture (MS): Course introduction	1. Book review: - Submit Top 3 Books	1. Davenport & Patil (2012) 2. Stair & Reynolds (2012) - CH 1,3 3. Pritzker, P., and May, W. (2015) - CH 2, App. A 4. Chapman et al. (2000) - CH 1,2 5. Spruit & Lytras (2018)	N/A
week 2	2018-09-10		Required attendance: - Tutorial 1A: Azure VMs - Tutorial 1B: Linux Bash	Regular lecture (MS): Big Data Engineering with Hadoop	Regular lecture (MS): Hadoop MapReduce and HDFS	1. Complete Tutorial 1: - Bash in Ubuntu 2. Start Tutorial 2: - Wordcount in Hadoop	- White (2015) - CH 1,2 - Dean & Ghemawat (2008) - Get & Read your selected book to review	LabA: Ubuntu 18 1 core, 2 MB
week 3	2018-09-17		Hands-on Tutorial 2: Wordcount in Hadoop	Required: Guest lecture (DV): UMCU/Neonatology on Big Data for Small Babies	Regular lecture (MS): Spark Architecture & Transformations	1. Complete Tutorial 2: - Wordcount in Hadoop 2. Start Tutorial 3: - Neonatology Part I	- White (2015) - CH 3,19 - Complete reading your selected book to review	LabB: Ubuntu 18 + Hadoop artifact 2 cores, 4 MB
week 4	2018-09-24	Python tutorial & Q/A	1. Complete Tutorial 3: - Neonatology Part I 2. Start Tutorial 3: - Neonatology Part II	Guest lecture (SM): ORTEC on Big Data Knowledge Discovery	Wrap-up lecture (MS): - Wide Transformations - Walkthrough of Tutorial 3 in Hadoop & Spark - Midterm Q/A	1. Submit Book 2-pager 2. Complete Tutorial 3: - Neonatology Parts I,II	- Complete readings above - Review tutorials to understand at the command level - Review _all_ lecture slides	LabB: Ubuntu 18 + Hadoop artifact 2 cores, 4 MB
week 5	2018-10-01	MIDTERM EXAM	NO LAB	Student pitch session: - The TOP-20 Data Science & Society books	Regular lecture (MB): Methods & Statistics I	1. Complete Tutorial 4: - Neonatology in Spark 2. Start Tutorial 5: - Statistics in Jupyter	- Lazer et al. (2014, March 28) - Broniatowski et al. (2014, July 28) - Chambers & Zaharia (2018) - CH 1	LabA: Ubuntu 18 DSVM 2 cores, 4 MB
week 6 (42)	2018-10-08		1. Complete Tutorial 5 - Statistics in Jupyter	Regular lecture (MB): - Methods & Statistics II	Required: Guest lecture (COM): UMCU/Epidemiology on Menopause and Cardiometabolic Disease Risk	1. Continue Tutorial 5: - Statistics in Jupyter 2. Start Tutorial 6: - Epidemiology Analytics in Jupyter Part I	- Chambers & Zaharia (2018) - CH 2	LabA: Ubuntu 18 DSVM 2 cores, 4 MB
week 7	2018-10-15		1. Explain assignment Big Spatial Data 2. Complete Tutorial 6 Part I - Epidemiology Analytics in Jupyter Part I	Regular lecture (MB+MS): - Methods & Statistics Wrap-up - SQL in Spark	Guest lecture (TB,JWvE): ESRI NL on Big Data in Geographical Information Systems	1. Complete Tutorial 6 Part I - Epidemiology Data Preprocessing 2. Start Tutorial 6 Part II - Epidemiology Analytics	- Chambers & Zaharia (2018) - CH 10	LabA: Ubuntu 18 DSVM 2 cores, 4 MB
week 8	2018-10-22		1. Complete Tutorial 6 Part II - Epidemiology Analytics	Guest lecture (WO): - CoreLifeAnalytics on Machine Learning in Cell Screening	Guest lecture (MM): UMCU/Julius on Big Data Ethics in Research, Privacy and Data Protection	1. Complete Tutorial 6 Part II - Epidemiology Data Analytics	- Complete ALL readings above - Review ALL tutorials to understand at the command level - Review ALL lecture slides	LabB: HDInsight VMs on Azure DataBricks cluster
week 9	2018-10-29		1. Complete Tutorial 6 Part II - Epidemiology Data Analytics 2. Work on Tutorial 6 Part III - Epidemiology Advanced Analytics	Guest lecture (FS,VM): UMCU/Psychiatry on Big Data in Psychiatry	FINAL LECTURE (MS): - "Towards Tomorrow: Trends in Data Science & Society" - Endterm Q/A	1. Complete Tutorial 6 Part III - Epidemiology Advanced Analytics	- Complete ALL readings above - Review ALL tutorials to understand at the command level - Review ALL lecture slides	LabB: HDInsight VMs on Azure DataBricks cluster
week 10	2018-11-05				ENDTERM EXAM			
week 1	2010-01-03				2ND CHANCE EXAM			

All multiple choice



Learning objectives in midterm exam

	Litera-ture	Work-shops	Guest talks	Regular lectures	TOTAL
Understand the role of data science and its societal impact	9		1		10
Recognise the knowledge discovery processes in applied data science	6				6
Identify trends and developments in big data engineering & analytics			1	15	16
Apply selected big data technologies to solve real-world problems		37		26	63
TOTAL	15	37	2	41	95



Extra assignment Big Spatial Data

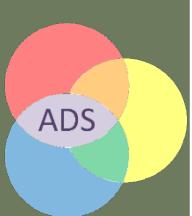
- › Playful Preparation for Guest lecture:
 - ESRI NL on Big Data in Geographical Information Systems
 - PDF available on Teams: [url](#)

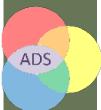
Materials from Dr. Sergio España

SQL in context and examples

*Refreshing your memory
& CH3 of Fundamentals of IS (Stair & Reynolds)*

Structured Query Language (SQL)'s DDL & DML



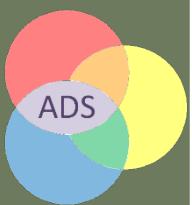


Structured Query Language

- › Abbreviation: SQL
- › Pronounced as: ‘S-Q-L’ [‘Ess-Que-El’]
- › Initially introduced as ‘SEQUEL’ in the 70s
- › SQL-92 is a 1992 ANSI standard
- › SQL:2011 is the current standard
- › Apache Spark implements subset of [SQL:2003](#).
- › “*Structured Query Language* is a domain-specific language for expressing relational operations over data” (CH 10)

SQL vs NoSQL

Quotes from: *Cattell, R. (2011). Scalable SQL and NoSQL data stores. ACM Sigmod Record, 39(4), 12-27.*





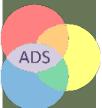
SQL vs NoSQL: *the survey*

- › Cattell, R. (2011). Scalable SQL and NoSQL data stores. *ACM Sigmod Record*, 39(4), 12-27.
- › *NoSQL systems generally have six key features:*
 1. the ability to horizontally scale “simple operation” throughput over many servers,
 2. the ability to replicate and to distribute (partition) data over many servers,
 3. a simple call level interface or protocol (in contrast to a SQL binding),
 4. a weaker concurrency model than the ACID transactions of most relational (SQL) database systems,
 5. efficient use of distributed indexes and RAM for data storage, and
 6. the ability to dynamically add new attributes to data records.



SQL vs NoSQL: *the survey*

- › A key feature of NoSQL systems is “shared nothing” horizontal scaling – replicating and partitioning data over many servers
- › No ACID transactional properties: updates are eventually propagated, but there are limited guarantees on the consistency of reads. Some authors suggest a “BASE” acronym in contrast to the “ACID” acronym:
 - **BASE** = Basically Available, Soft state, Eventually consistent
 - **ACID** = Atomicity, Consistency, Isolation, and Durability
- › The idea is that by giving up ACID constraints, one can achieve much higher performance and scalability.
- › Sounds a lot like... Hadoop??



SQL vs NoSQL: The argument for relational

1. If new relational systems can do everything that a NoSQL system can, with analogous performance and scalability, and with the convenience of transactions and SQL, why would you choose a NoSQL system?
2. Relational DBMSs have taken and retained majority market share over other competitors in the past 30 years: network, object, and XML DBMSs.
3. Successful relational DBMSs have been built to handle other specific application loads in the past: read-only or read-mostly data warehousing, OLTP on multi-core multi-disk CPUs, in-memory databases, distributed databases, and now horizontally scaled databases.
4. While we don't see "one size fits all" in the SQL products themselves, we do see a common interface with SQL, transactions, and relational schema that give advantages in training, continuity, and data interchange.

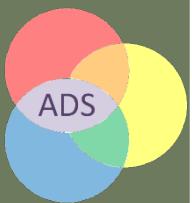


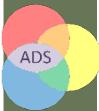
SQL vs NoSQL: The counter-argument for NoSQL

1. We haven't yet seen good benchmarks showing that RDBMSs can achieve scaling comparable with NoSQL systems like Google's BigTable.
2. If you only require a lookup of objects based on a single key, then a key-value store is adequate and probably easier to understand than a relational DBMS. Likewise for a document store on a simple application: you only pay the learning curve for the level of complexity you require.
3. Some applications require a flexible schema, allowing each object in a collection to have different attributes. While some RDBMSs allow efficient “packing” of tuples with missing attributes, and some allow adding new attributes at runtime, this is uncommon.
4. A relational DBMS makes “expensive” (multinode multi-table) operations “too easy”. NoSQL systems make them impossible or obviously expensive for programmers. [un-ease of use??]
5. While RDBMSs have maintained majority market share over the years, other products have established smaller but non-trivial markets in areas where there is a need for particular capabilities, e.g. indexed objects with products like BerkeleyDB, or graph-following operations with object-oriented DBMSs.

SQL DDL

Data Definition Language





SQL data definitions

- › **CREATE**
 - Creation of database objects
- › **ALTER**
 - Modify the structure and/or the characteristics of database objects
- › **DROP**
 - Deletion of database objects
- › **TRUNCATE**
 - Deletion of data in tables without altering the structure



SQL DDL and SQL DML

- › SQL is no ‘classical’ programming language, it is more a ‘data sub language’
- › Core parts of SQL:
 - **Data definition language (DDL)**: used to define database structures
 - **Data manipulation language (DML)**: define, update, and request data (queries)



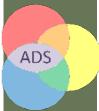
The CREATE TABLE statement

- › Creation of tables
 - The SQL CREATE TABLE statement

- **CREATE TABLE EMPLOYEE(**

EmpID	Integer	PRIMARY KEY,
EmpName	Char(25)	NOT NULL
);		





Composite PRIMARY KEYS

- › Creation of database tables with composite keys

– CREATE TABLE EMP_SKILL(

EmplD	Integer	NOT NULL,
SkillID	Integer	NOT NULL,
SkillLevel	Integer	NULL,
CONSTRAINT EmpSkill_PK		PRIMARY KEY (EmplD, SkillID)
);		

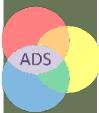
Emp_Skill		
PK	EmplD	Integer
PK	SkillID	Integer
	SkillLevel	Integer



The FOREIGN KEY constraint (1/2)

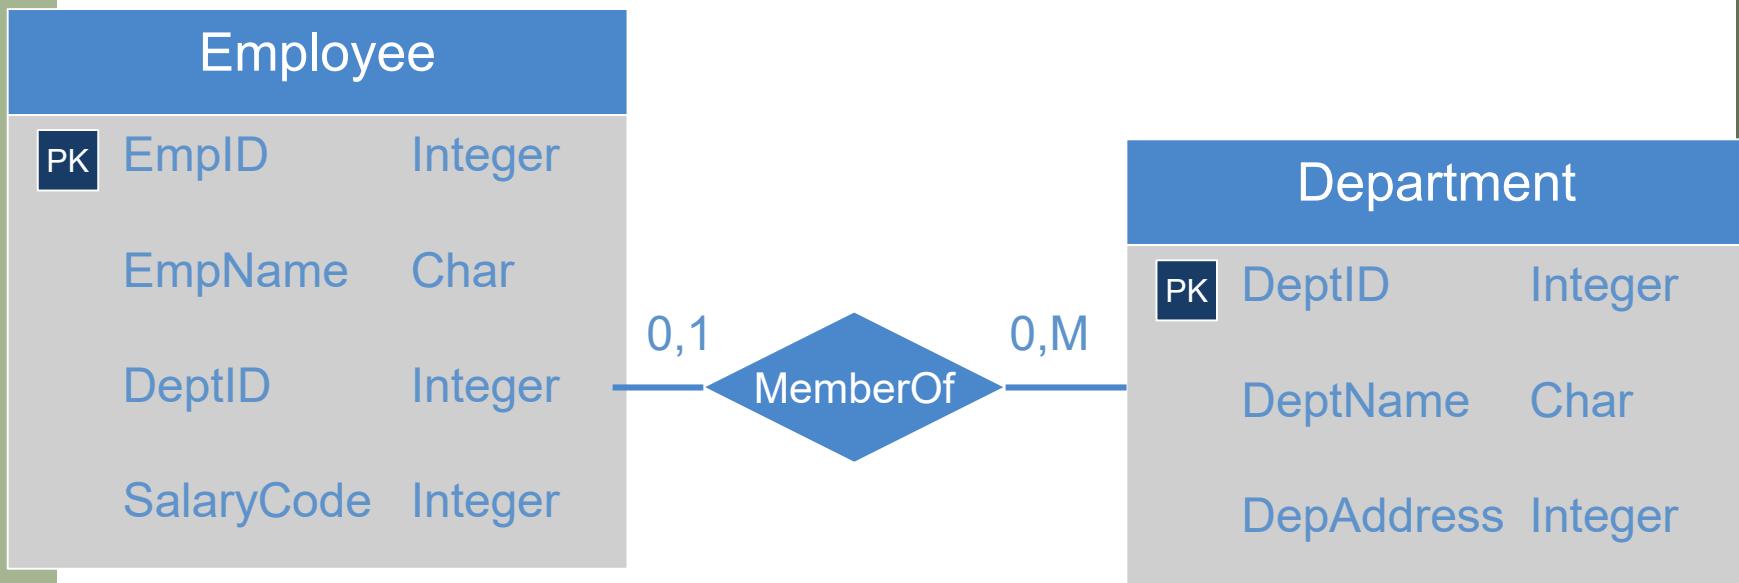
- › Creation of database tables using PRIMARY KEY and FOREIGN KEY constraints

```
– CREATE TABLE EMP_SKILL(  
    EmpID      Integer      NOT NULL,  
    SkillID     Integer      NOT NULL,  
    SkillLevel   Integer      NULL,  
    CONSTRAINT EmpSkill_PK PRIMARY KEY (EmpID, SkillID),  
    CONSTRAINT Emp_FK      FOREIGN KEY(EmpID)  
                           REFERENCES EMPLOYEE(EmpID),  
    CONSTRAINT Skill_FK    FOREIGN KEY(SkillID)  
                           REFERENCES SKILL(SkillID)  
);
```



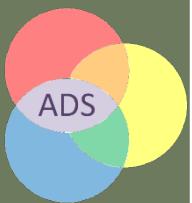
Deletion of database objects

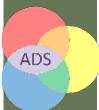
- › DROP: Deletion of unwanted database objects
- › Warning! The DROP keyword permanently deletes the data object along with the data related to that object
 - › `DROP TABLE EMPLOYEE;`



SQL DML

Data Manipulation Language





The INSERT keyword

- › Use the INSERT keyword to add a row in a table
- › Non-numerical data needs to be within single quotes ('')
- › `INSERT INTO EMPLOYEE VALUES(91, 'Smither', 12);`
- › `INSERT INTO EMPLOYEE (EmpID, SalaryCode) VALUES (62, 11);`

Employee			
EmpID	EmpName	DeptID	SalaryCode
91	Smither	12	
62	Null!	Null!	11



The UPDATE keyword

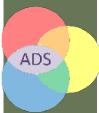
- › Modify values in an existing row (or a collection of rows)

```
UPDATE EMPLOYEE  
SET     Phone '791-555-123'  
WHERE   EmpID = 29;
```

```
UPDATE EMPLOYEE  
SET     DeptID = 14  
WHERE   EmpName LIKE 'Kr%';
```

Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	12	11	445-445-334
29	Homer	15	8	666-666-666



The DELETE keyword

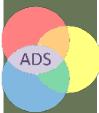
- › Delete a row or a set of rows

```
DELETE FROM EMPLOYEE  
WHERE EmpID = 29;
```

```
DELETE FROM EMPLOYEE  
WHERE EmpName LIKE 'Kr%';
```

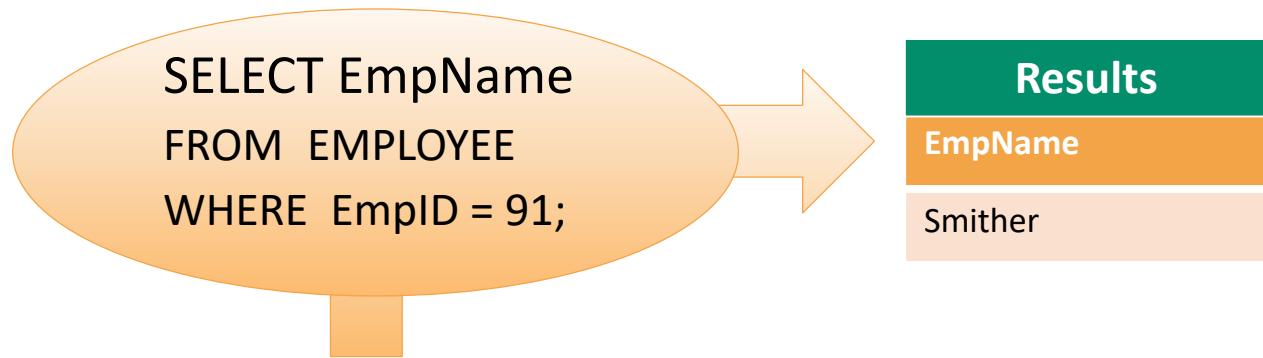
Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



Queries

- › A query retrieves data from one or more tables and creates a new (temporary) table
- › **SELECT** is the best-known SQL statement



Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



Showing all values from a table

- › Use an asterisk (*) to show all values in a table that match certain criteria.

Results

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333

```
SELECT *
FROM EMPLOYEE
WHERE EmpID = 91;
```

Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



The DISTINCT keyword

- › The DISTINCT keyword can be added to the SELECT statement to prevent that duplicate rows are shown

```
SELECT DISTINCT DeptID  
FROM EMPLOYEE;
```

Results
DeptID
12
14
15

Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



The WHERE keyword (1/2)

- With the WHERE keyword criteria can be provided which the shown records should meet

```
SELECT EmpName  
FROM EMPLOYEE  
WHERE DeptID = 15;
```

Results
EmpName
Homer
Krall

Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



The WHERE keyword (2/2)

› The WHERE clause may include the following symbols:

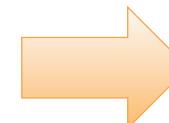
- Equals “=”
- Not Equals “<>”
- Greater than “>”
- Less than “<”
- Greater than or Equal to “>=”
- Less than or Equal to “<=”



Logical AND- and OR-operators

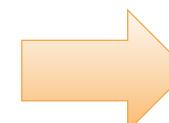
- › Multiple matching criteria can be specified by using:
 - AND: represents an intersection of data collections
 - OR: represents a union of data collections

```
SELECT      EmpName  
FROM        EMPLOYEE  
WHERE       DeptID < 13 OR  
                  DeptID > 14;
```



Results
EmpName
Smither
Homer
Krall

```
SELECT      EmpName  
FROM        EMPLOYEE  
WHERE       DeptID = 15 AND  
                  SalaryCode <= 8;
```



Results
Homer



The IN keyword

- › The WHERE clause can include the IN keyword to make sure that certain column values are included in the query result

```
SELECT EmpName  
FROM EMPLOYEE  
WHERE DeptID IN (4, 12, 14);
```

Results
EmpName
Smither
Krusty

Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



Logical NOT-operator

- › The logical NOT-operator is used to prevent records from being part of the resulting set of records

```
SELECT EmpName  
FROM EMPLOYEE  
WHERE DeptID NOT IN (4, 12, 14);
```

Results
EmpName
Homer
Krall

Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



The BETWEEN keyword

- › The BETWEEN keyword is used to show records where the values in a column are between a minimum and maximum value

```
SELECT EmpName  
FROM EMPLOYEE  
WHERE SalaryCode BETWEEN 7 AND 9;
```

Results
EmpName
Homer

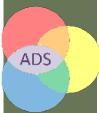
Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



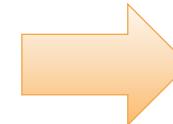
The LIKE keyword

- › By means of the LIKE keyword partially complete values can be searched for
- › Use wildcards to find records that are equal to a certain string value
 - Percentage symbol (%) for multiple arbitrary characters
 - Underscore (_) for a single arbitrary character



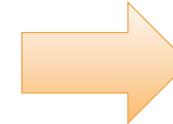
Examples

```
SELECT EmpID  
FROM EMPLOYEE  
WHERE EmpName LIKE 'Kr%';
```



Results
EmpID
33
32

```
SELECT EmpID  
FROM EMPLOYEE  
WHERE Phone LIKE '___-__-33_';
```



Results
EmpID
91
33

Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



Sorting of results

- › Query results are sorted by means of the ORDER BY clause:

```
SELECT EmpID, EmpName  
FROM EMPLOYEE  
ORDER BY EmpName;
```

Results	
EmpID	EmpName
29	Homer
32	Krall
33	Krusty
91	Smither

Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



Built-in SQL functions

› SQL offers various built-in functions:

- COUNT : counts the number of rows that match the criteria as posed in the query
- MIN: search for the minimum value in a certain column
- MAX: search for the maximum value in a certain column
- SUM: computes the sum of values in a certain column
- AVG: computes the average of the values in a certain column



Examples

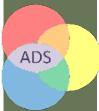
```
SELECT COUNT(DeptID)  
FROM EMPLOYEE;
```

Results
COUNT(DeptID)
4

```
SELECT COUNT(DISTINCT DeptID)  
FROM EMPLOYEE;
```

Results
COUNT(DISTINCT DeptID)
3

Employee				
EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



Examples

```
SELECT MIN(Hours) AS MinimumHours,  
       MAX(Hours) AS MaximumHours,  
       AVG(Hours) AS AverageHours  
FROM    PROJECT  
WHERE ProjID > 7;
```

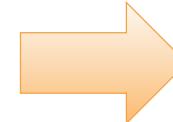
```
SELECT SUM(Hours) AS TotalHours,  
      FROM  PROJECT;
```



Showing sub totals: GROUP BY

- Sub totals are computed by means of the GROUP BY clause

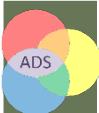
```
SELECT DeptID,  
       COUNT(*) AS NumEmp  
  FROM EMPLOYEE  
 GROUP BY DeptID;
```



Results	
DeptID	NumEmp
12	1
14	1
15	2

Employee

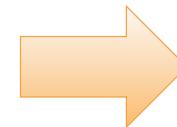
EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



Showing sub totals: GROUP BY

- › Sub totals are computed by means of the GROUP BY clause
- › The HAVING clause is used to limit how much data is shown

```
SELECT DeptID,  
       COUNT(*) AS NumEmp  
  FROM EMPLOYEE  
 GROUP BY DeptID  
 HAVING COUNT(*) > 1;
```



Results	
DeptID	NumEmp
15	2

Employee

EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



Example: “HAVING COUNT(*) > 1” for Privacy

- › Example of people with a disease, stored with their postal code and their ethnicity.
- › Implements “2-anonymity” when the person is not allowed to see detailed results (SELECT * FROM XXX WHERE allowed = 0)
- › And all the detailed results when the person is allowed (SELECT * FROM XXX WHERE allowed = 1).
- › Toledo,C. van, & Spruit,M. (2016). Adopting privacy regulations in a data warehouse: A case of the anonymity versus utility dilemma. In Fred,A. et al. (Ed.), *Proceedings of the 8th International Joint Conference on Knowledge Discovery, Knowledge Engineering and Knowledge Management* (pp. 67–72). KDIR 2016, November 11-13, 2016, Porto, Portugal: ScitePress. [[pdf](#)]

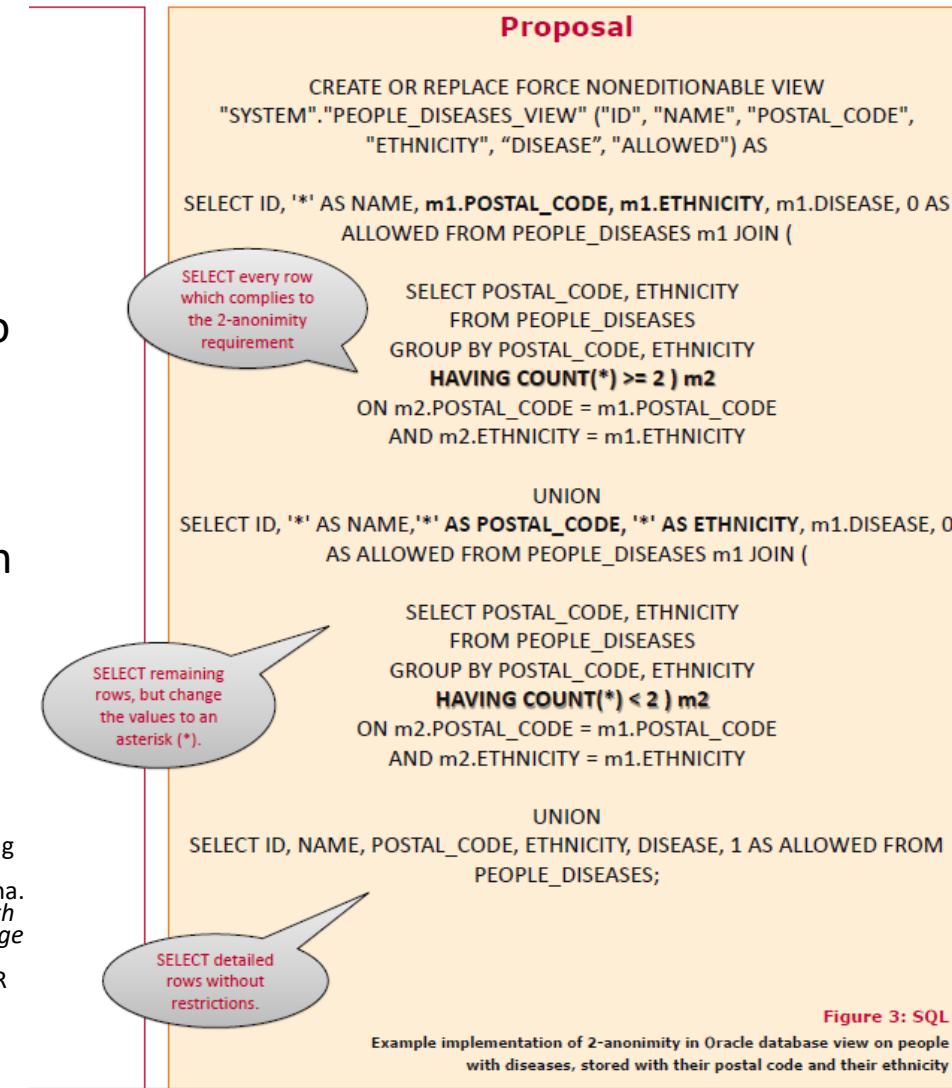


Figure 3: SQL

Example implementation of 2-anonymity in Oracle database view on people with diseases, stored with their postal code and their ethnicity



Data from multiple tables

› Subqueries

- The result of a query is a subset of the data
- Therefore, a query can be used as input for another query
- This is called a subquery

› Joins

- A different way to combine data can be realized with a join
 - › Join (can also be called Inner Join)
 - › Left Outer Join
 - › Right Outer Join



Subquery example

```
SELECT EmpName  
FROM EMPLOYEE  
WHERE DeptID IN  
(SELECT DeptID  
FROM DEPARTMENT  
WHERE DeptName  
LIKE 'Accou%');
```



Results	
EmpName	
Smither	
Krusty	

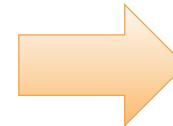
Department		
DeptID	DeptName	DeptAddress
12	Accountancy	Floor 3
14	Accounts VIP	Floor 4
15	Production	Warehouse A

Employee				
EmpID	EmpName	DeptID	SalaryCode	Phone
91	Smither	12	10	773-334-333
33	Krusty	14	11	445-445-334
29	Homer	15	8	791-555-123
32	Krall	15	10	554-678-889



Join example

```
SELECT EmpName  
FROM EMPLOYEE AS E,  
      DEPARTMENT AS D  
WHERE E.DeptID = D.DeptID  
      AND D.DeptName  
      LIKE 'Accou%';
```

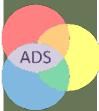


Results	
EmpName	
Smither	
Krusty	

Department		
DeptID	DeptName	DeptAddress
12	Accountancy	Floor 3
14	Accounts VIP	Floor 4

Employee JOIN Department

E.EmpID	E.EmpName	E.DeptID	E.SalaryCode	E.Phone	D.DeptID	D.DeptName	D.DeptAddress
91	Smither	12	10	773-334-333	12	Accountancy	Floor 3
33	Krusty	14	11	445-445-334	14	Accounts VIP	Floor 4
29	Homer	15	8	791-555-123	15	Production	Warehouse A
32	Krall	15	10	554-678-889	15	Production	Warehouse A
29	Homer	15	8	791-555-123	10	554-678-889	
32	Krall	15	10	554-678-889	10	554-678-889	



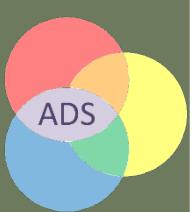
The JOIN...ON syntax

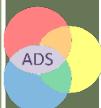
- › Moving the JOIN syntax to the FROM clause:

```
SELECT EmpName  
FROM EMPLOYEE AS E JOIN DEPARTMENT AS D  
      ON E.DeptID = D.DeptID  
WHERE D.DeptName LIKE 'Account%';
```

SQL in Apache Spark

Best of Both Worlds!





SQL in the Spark Architecture

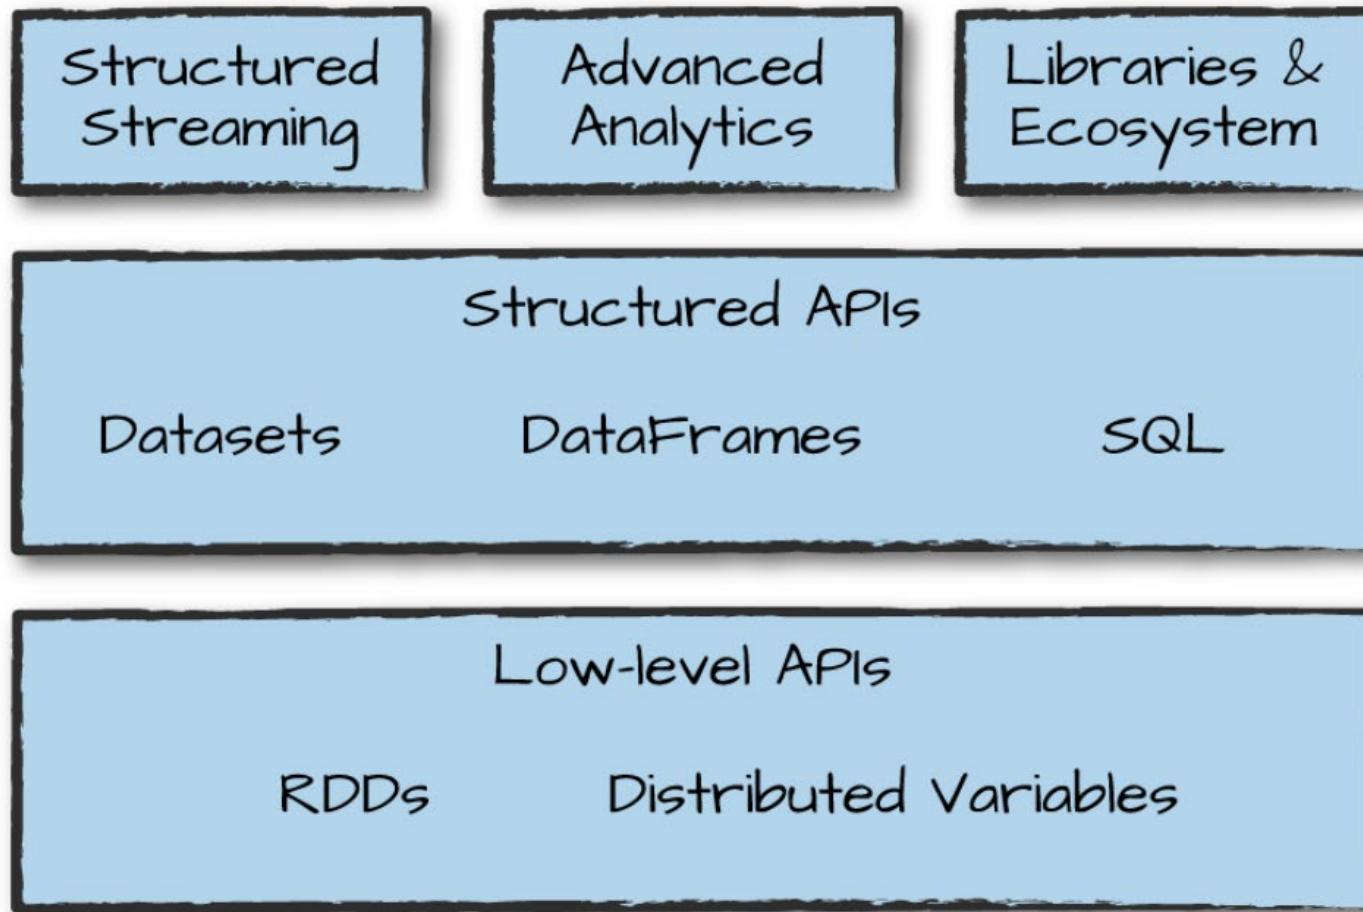
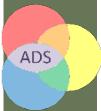
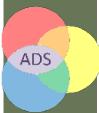


Figure 1-1. Spark's toolkit



Big Data and SQL: How, until Spark 2?

- › Apache Hive
 - Before Spark's rise, Hive was the de facto big data SQL access layer
 - Originally developed at Facebook, Hive became an incredibly popular tool across industry for performing SQL operations on big data
 - Why important for Hadoop adoption and cross-industry success?
 - › because analysts could run SQL queries
- › Why are now many users using Spark SQL, instead of...?
 - Resilient Distributed Datasets (RDDs)
- › Compatible with Hive
- › Interoperable with DataFrames



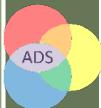
Using SQL in Spark

- › **Spark SQL CLI:** ./bin/spark-sql
- › **Spark's SQL API:**
 - via the method sql on the SparkSession object.
 - returns a DataFrame
 - spark.sql("SELECT 1 + 1").show()
 - › Is a *lazy* transformation
- › A Table is logically equivalent to a DataFrame

```
CREATE TABLE flights_csv (
    DEST_COUNTRY_NAME STRING,
    ORIGIN_COUNTRY_NAME STRING COMMENT "remember, the US will be most prevalent",
    count LONG)
USING csv OPTIONS (header true, path '/data/flight-data/csv/2015-summary.csv')

CREATE TABLE flights (
    DEST_COUNTRY_NAME STRING, ORIGIN_COUNTRY_NAME STRING, count LONG)
USING JSON OPTIONS (path '/data/flight-data/json/2015-summary.json')

INSERT INTO flights_from_select
    SELECT DEST_COUNTRY_NAME, ORIGIN_COUNTRY_NAME, count FROM flights LIMIT 20
```



SQL Queries in Spark

Select Statements

Queries in Spark support the following ANSI SQL requirements (here we list the layout of the SELECT expression):

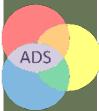
```
SELECT [ALL|DISTINCT] named_expression[, named_expression, ...]
      FROM relation[, relation, ...]
      [lateral_view[, lateral_view, ...]]
      [WHERE boolean_expression]
      [aggregation [HAVING boolean_expression]]
      [ORDER BY sort_expressions]
      [CLUSTER BY expressions]
      [DISTRIBUTE BY expressions]
      [SORT BY sort_expressions]
      [WINDOW named_window[, WINDOW named_window, ...]]
      [LIMIT num_rows]

named_expression:
  : expression [AS alias]

relation:
  | join_relation
  | (table_name|query|relation) [sample] [AS alias]
  : VALUES_(expressions)[, (expressions), ...]
    [AS (column_name[, column_name, ...])]

expressions:
  : expression[, expression, ...]

sort_expressions:
  : expression [ASC|DESC][, expression [ASC|DESC], ...]
```



A different IF-THEN

› **case...when...then** Statements

- conditionally replace values in SQL queries
- essentially the equivalent of programmatic if statements

```
SELECT
    CASE WHEN DEST_COUNTRY_NAME = 'UNITED STATES' THEN 1
          WHEN DEST_COUNTRY_NAME = 'Egypt' THEN 0
          ELSE -1 END
FROM partitioned_flights
```