

ML for Human Vision and Language

MSc Artificial Intelligence

Lecturer: Tejaswini Deoskar
t.deoskar@uu.nl

UiL-OTS, Utrecht University

Block 1, 2019

Parts-of-speech

word classes, lexical classes, grammatical classes, lexical tags, ...

POS tagging (inference): *to determine the POS tag for a particular instance (token) of a word in context.*

EX	VBD	JJ	NN	IN	DT	NN
There	was	still	lemonade	in	the	bottle

- Various granularities: fine-grained tags/coarse-grained tags

POS tagging: Why do we care?

- first step towards syntactic analysis; simpler and faster than full syntactic parsing
- POS tagging task also helps introduce useful techniques: **Hidden Markov models**(HMMs) or **Recurrent Neural networks** (RNNs), which are used for many other sequence labelling or sequence modelling tasks.
- A variety of (AI-related) applications are sequence labelling/sequence modelling tasks : robot actions, music cognition/generation, etc.
- POS tags can be useful features in e.g. text classification, authorship identification, text-to-speech etc.

Other tagging or sequence labelling tasks

- **Named Entity Recognition (NER)** e.g., label words as belonging to persons (PER), organizations (ORG), locations(LOC), or none of the above:
 - * Donald/PER Trump/PER tweeted/NON from/NON the/NON White/LOC House/LOC today/NON ./NON

Other tagging or sequence labelling tasks

- **Named Entity Recognition (NER)** e.g., label words as belonging to persons (PER), organizations (ORG), locations(LOC), or none of the above:
 - * Donald/PER Trump/PER tweeted/NON from/NON the/NON White/LOC House/LOC today/NON ./NON
- **Information field segmentation:** Given semi-structured text (classified advertisement, bibliography entry), identify which words belong to which fields (price/ size/ location, author /title/ year)
 - * 3BR/SIZE flat/TYPE in/NON Utrecht/LOC ,/NON near/LOC centraal/LOC station/LOC ./NON Bright/FEAT ,/NON well/FEAT maintained/FEAT ...

Problems with bi-gram (N-gram) tagging

Bigram: The still smoking remains of the campfire
 DT

Problems with bi-gram (N-gram) tagging

Bigram:	The	still	smoking	remains	of	the	campfire
	DT	JJ					

Problems with bi-gram (N-gram) tagging

	The	still	smoking	remains	of	the	campfire
Bigram:	DT	JJ	NN	VBZ	...		

Problems with bi-gram (N-gram) tagging

	The	still	smoking	remains	of	the	campfire
Bigram:	DT	JJ	NN	VBZ	...		
Intended:	DT	RB	VBG	NNS	IN	DT	NN

Problems with bi-gram (N-gram) tagging

	The	still	smoking	remains	of	the	campfire
Bigram:	DT	JJ	NN	VBZ	...		
Intended:	DT	RB	VBG	NNS	IN	DT	NN

- One incorrect tagging choice might have unintended effects:
- No lookahead: choosing the “most probable” tag at one stage might lead to highly improbable choice later.

Problems with bi-gram (N-gram) tagging

	The	still	smoking	remains	of	the	campfire
Bigram:	DT	JJ	NN	VBZ	...		
Intended:	DT	RB	VBG	NNS	IN	DT	NN

- One incorrect tagging choice might have unintended effects:
- No lookahead: choosing the “most probable” tag at one stage might lead to highly improbable choice later.
- We want to find the **overall most likely** tagging sequence, given the bigram (n-gram) frequencies. This is what the Hidden Markov Model (HMM) approach achieves.

Finding the most likely sequence

- Formalising the model
- Expressing in terms of Hidden Markov Model (HMM)
- Viterbi algorithm to find the best (tag) sequence
- Other sequence models

Stochastic tagging: Formalizing the problem

A sentence is a sequence of n words $(w_1 w_2 \dots w_n)$: w_1^n

POS tags for the sentence is a sequence of n tags $(t_1 t_2 \dots t_n)$: t_1^n

Stochastic tagging: Formalizing the problem

A sentence is a sequence of n words ($w_1 w_2 \dots w_n$): w_1^n

POS tags for the sentence is a sequence of n tags ($t_1 t_2 \dots t_n$): t_1^n

We want to find the most probable tag sequence t_1^n , given the word sequence w_1^n

$$\hat{t}_1^n = \arg \max_{t_1^n} P(t_1^n | w_1^n)$$

Using Bayes' Rule

$$\hat{t}_1^n = \arg \max_{t_1^n} P(t_1^n | w_1^n) = \arg \max_{t_1^n} \frac{P(w_1^n | t_1^n) P(t_1^n)}{P(w_1^n)}$$

- Ignore the denominator, since best tag sequence does not depend on prob. of word sequence

$$\hat{t}_1^n = \arg \max_{t_1^n} P(w_1^n | t_1^n) P(t_1^n)$$

- $P(t_1^n)$: prior probability of the tag sequence
- $P(w_1^n | t_1^n)$: likelihood of data

Simplifying assumptions

- Estimating these probabilities from corpus data will be hard (sparsity)

$P(\text{"Time flies like an arrow"} | \text{NN VB IN DT NN})$

$P(\text{NN VB IN DT NN})$

- Make simplifying assumptions for each term

Simplifying assumptions (Independence assumptions)

For prior term :

$$P(t_1^n) = \prod_{i=1}^n P(t_i|t_{i-1})$$

- The probability of a tag is independent of all but the previous tag (bi-grams of tags)

Simplifying assumptions (Independence assumptions)

For prior term :

$$P(t_1^n) = \prod_{i=1}^n P(t_i|t_{i-1})$$

- The probability of a tag is independent of all but the previous tag (bi-grams of tags)

For likelihood term:

$$P(w_1^n|t_1^n) = \prod_{i=1}^n P(w_i|t_i)$$

- The probability of a word depends only on its POS tag
 - * Independent of other words in the sentence
 - * Independent of other POS tags

Bi-gram tagger

$$\hat{t}_1^n = \arg \max_{t_1^n} P(w_1^n | t_1^n) P(t_1^n)$$

$$\hat{t}_1^n \approx \arg \max_{t_1^n} \prod_{i=1}^n P(w_i | t_i) P(t_i | t_{i-1})$$

Estimating the parameters of the model from a corpus

Specific probabilities of $P(t_i|t_{i-1})$ and $P(w_i|t_i)$ are estimated from a tagged (supervised) corpus:

Estimating the parameters of the model from a corpus

Specific probabilities of $P(t_i|t_{i-1})$ and $P(w_i|t_i)$ are estimated from a tagged (supervised) corpus:

Tag Transition probabilities :

$$P(t_i|t_{i-1}) = \frac{\text{count}(t_{i-1}, t_i)}{\text{count}(t_{i-1})}$$

- A noun but not a verb is very likely to follow a determiner or adjective:

That/DT flight/NN

The/DT yellow/JJ hat/NN

- Example estimated from Brown corpus:

$$P(\text{NN} | \text{DT}) = \frac{\text{count}(\text{DT, NN})}{\text{count}(\text{DT})} = \frac{56,509}{116,454} = 0.49$$

Estimating the parameters of the model from a corpus

Word likelihood probabilities :

$$P(w_i|t_i) = \frac{\text{count}(t_i, w_i)}{\text{count}(t_i)}$$

- If the tag is VBZ (3 Person Singular Present Tense), then 'is' is very likely to be the verb.

¹Example from J&M 2 edition 5.5

Estimating the parameters of the model from a corpus

Word likelihood probabilities :

$$P(w_i|t_i) = \frac{\text{count}(t_i, w_i)}{\text{count}(t_i)}$$

- If the tag is VBZ (3 Person Singular Present Tense), then 'is' is very likely to be the verb.
- Example estimated from Brown corpus:

$$P(\text{is} | \text{VBZ}) = \frac{\text{count}(\text{VBZ}, \text{is})}{\text{count}(\text{VBZ})} = \frac{10,073}{21,627} = 0.47$$

1

¹Example from J&M 2 edition 5.5

What can we do with this model - I?

- Compute the probability of a tagged sentence, if we know the parameters.

$$P(w_1^n, t_1^n) = \prod_{i=1}^n P(t_i | t_{i-1}) P(w_i | t_i)$$

- Comparison to a Language Model:
 - * a LM gives us the probability of a sentence, whereas this model here gives us the probability of a **tagged sentence**.
- (Caveat: A start $< s >$ and stop $< /s >$ symbol is added to beginning and end of sentences.)

What can we do with this model - II?

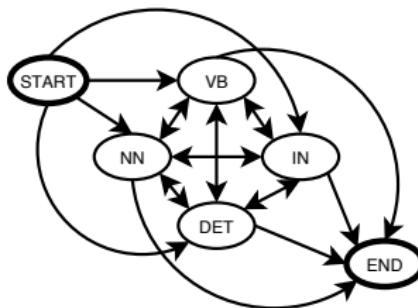
- Tagging: use the model to find the best tag sequence t_1^n for an untagged sentence w_1^n

$$\hat{t}_1^n = \arg \max_{t_1^n} P(t_1^n | w_1^n)$$

- * thus, **Hidden** Markov Model
- * **Markov** because of Markov assumption (tag only depends on immediately previous tag)
- * **Hidden** because we only observe the words/emissions; the tags/states are hidden (or **latent**) variables.

Probabilistic Finite State Machine

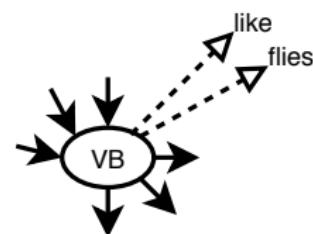
- Sentences are generated by walking through **states** in a graph.
Each state represents a tag.
- Each state emits an **output** (a word)



transition probability

Prob. of moving from state s to s' :

$$P(t_i = s' | t_{i-1} = s)$$



emission probability

Prob of emitting w from state s :

$$P(w_i = w | t_i = s)$$

Hidden Markov Model (HMM)

HMM is a very general model for sequences.

The elements of an HMM:

- a set of states (here: the **tags**)
- an output alphabet (here: **words**)
- initial state (here: beginning of sentence)
- state transition probabilities (here: $p(t_i | t_{i-1})$)
- symbol emission probabilities (here: $p(w_i | t_i)$)

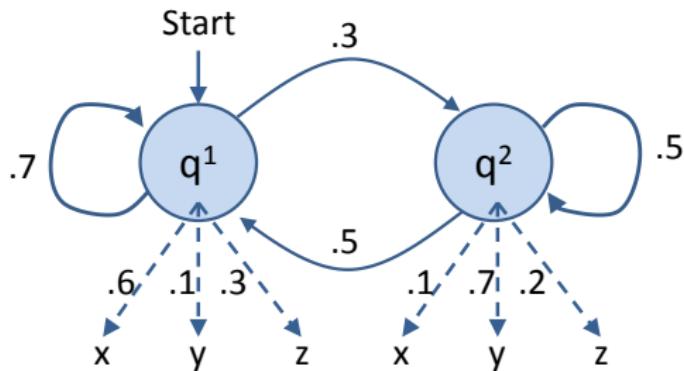
Notation

- Sequence of observations over time o_1, o_2, \dots, o_T
 - * here, **words** in sentence
- Vocabulary size V of possible observations
- Set of possible states q^1, q^2, \dots, q^N
 - * here, **tags** (*see note on notation on next slide*)
- A , an $N \times N$ matrix of transition probabilities
 - * a_{ij} : the prob. of transitioning from state i to j
- B , an $N \times V$ matrix of output probabilities
 - * $b_i(o_t)$: the prob. of emitting o_t from state i

Note on notation

- J&M use q_1, q_2, \dots, q_N for set of states, but also use q_1, q_2, \dots, q_T for state sequence over time
- We'll use q^i for state names, and q_t for state at time t
- So, we could have $q_t = q^i$, meaning : the state we're in at time t is q^i

HMM toy example with new notation



- States $\{q^1, q^2\}$ (or $\{s, q^1, q^2\}$)
- Output alphabet $\{x, y, z\}$

Adapted from Manning & Schuetze, Fig 9.2

Transition and Output Probabilities

- Transition matrix A : $a_{ij} = P(q^j|q^i)$

$< s >$	q^1	q^2
q^1	0.7	0.3
q^2	0.5	0.5

Transition and Output Probabilities

- Transition matrix A : $a_{ij} = P(q^j|q^i)$

$< s >$	q^1	q^2
q^1	0.7	0.3
q^2	0.5	0.5

- Output matrix B : $b_i(o) = P(o|q^i)$ for output o

	x	y	z
q^1	0.6	0.1	0.3
q^2	0.1	0.7	0.2

- $\lambda = (A, B)$ are the parameters of our HMM

Joint Probability of (states, outputs)

- Using the new notation, given state sequence $Q = (q_1 \dots q_T)$ and output sequence $O = (o_1 \dots o_T)$, we have:

$$P(O, Q | \lambda) = \prod_{t=1}^T P(o_t | q_t) \ P(q_t | q_{t-1})$$

$$P(O, Q | \lambda) = \prod_{t=1}^T b_{q_t}(o_t) \ a_{q_{t-1}q_t}$$

Search for the best tag sequence : Viterbi

- For any specific tag sequence t_1^n , it is easy to calculate

$$P(w_1^n, t_1^n) = \prod_{i=1}^n P(t_i | t_{i-1}) P(w_i | t_i)$$

- So why can't we enumerate all possible t_1^n sequences, compute their probability, and choose the best one?

The	still	smoking	remains
DT	RB	NN	NNS
JJ	VBG		VBZ
:	:		:

Enumeration won't work

- Suppose we have c possible tags for each of the n words in a sentence?
- How many possible tag sequences?

Enumeration won't work

- Suppose we have c possible tags for each of the n words in a sentence?
- How many possible tag sequences?
- There are c^n possible tag sequences: the number grows exponentially in the length n
- For all but sentences of small length, too many sequences to enumerate: $c = 10$

10 word sent. $\Rightarrow 10,000,000,000$ tag sequences

15 word sent. $\Rightarrow 1,000,000,000,000,000$ tag sequences

Overview: HMM for sequence tagging

- Given a sentence $O = o_1 \dots o_T$, with tags $Q = q_1 \dots q_T$, we want to compute $P(O, Q)$ as

$$P(O, Q) = \prod_{t=1}^T P(o_t | q_t) P(q_t | q_{t-1})$$

$$P(w_1^n, t_1^n) = \prod_{i=1}^n P(w_i | t_i) P(t_i | t_{i-1})$$

- But we want to find $\arg \max_Q P(Q|O)$ without enumerating all possible Q
 - * Use Viterbi algorithm to store partial computations

Finding the best sequence : Viterbi

- The **Viterbi** algorithm finds the best sequence without explicitly enumerating all sequences.
- Example of a **dynamic programming** or **memoisation** algorithm.
- Efficient, **exact** inference.
- stores partial results in a **chart** to avoid recomputing them.

Viterbi High level picture

Viterbi intuition: the best path of length t ending in state q must include the best path of length $t - 1$ to the previous state. (t now a time step, not a tag)

Viterbi High level picture

Viterbi intuition: the best path of length t ending in state q **must** include the best path of length $t - 1$ to the previous state. (t now a time step, not a tag)

- thus, we do not have to calculate the full best path at every time step!

High level picture:

- Find the best path of length $t - 1$ to each state.
- Consider extending each of those by 1 step, to state q
- Take the best of those options as the best path to state q

Viterbi Intuition

$< s >$	one	dog	bit	$< /s >$
$< s >$	CD	NN	NN	$< /s >$
	NN	VB	VBD	
	PRP			

Viterbi Intuition

$< s >$	one	dog	bit	$< /s >$
$< s >$	CD	NN	NN	$< /s >$
	NN	VB	VBD	
	PRP			

- Suppose we have already computed
 1. The best tag sequence for $< s > \dots$ bit that ends in NN.
 2. The best tag sequence for $< s > \dots$ bit that ends in VBD.

Viterbi Intuition

$< s >$	one	dog	bit	$< /s >$
$< s >$	CD	NN	NN	$< /s >$
	NN	VB	VBD	
	PRP			

- Suppose we have already computed
 1. The best tag sequence for $< s > \dots$ bit that ends in NN.
 2. The best tag sequence for $< s > \dots$ bit that ends in VBD.
- Then, the best **full** sequence would be either
 - * sequence (1) extended to include $< /s >$, or
 - * sequence (2) extended to include $< /s >$.

Viterbi Intuition

$< s >$	one	dog	bit	$< /s >$
$< s >$	CD	NN	NN	$< /s >$
	NN	VB	VBD	
	PRP			

- Suppose we have already computed
 1. The best tag sequence for $< s > \dots \text{bit}$ that ends in NN.
 2. The best tag sequence for $< s > \dots \text{bit}$ that ends in VBD.
- Then, the best **full** sequence would be either
 - * sequence (1) extended to include $< /s >$, or
 - * sequence (2) extended to include $< /s >$.
- Similarly, to get the best tag sequence for $< s > \dots \text{bit}$ that ends in NN, we could extend one of
 1. The best tag sequence for $< s > \dots \text{dog}$ that ends in NN.
 2. The best tag sequence for $< s > \dots \text{dog}$ that ends in VB.

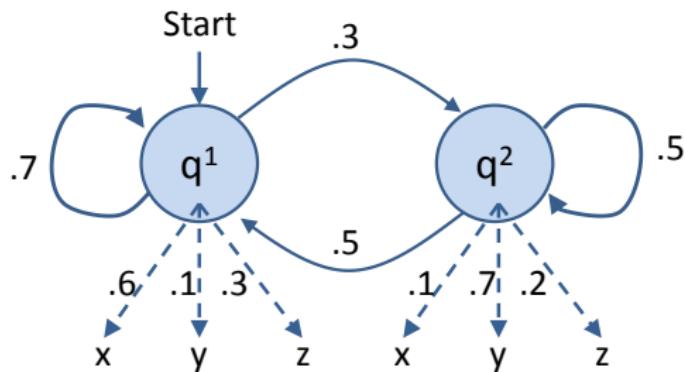
Viterbi Algorithm

- Use a chart $N \times T$ to store partial results as we go
 - * $v(j, t)$ is the probability of the best state sequence for $o_1 \dots o_t$ that ends in state j
- Fill in columns from left to right, with

$$v(j, t) = \max_{i=1}^N v(i, t-1) \cdot a_{ij} \cdot b_j(o_t)$$

- Store a backtrace to show, for each cell, which state at $t-1$ we came from.

Toy example



- States $\{q^1, q^2\}$ (or $\{\langle s \rangle, q^1, q^2\}$)
- Output alphabet $\{x, y, z\}$

Adapted from Manning & Schuetze, Fig 9.2

Toy example

- Suppose $O = x z y$.
- Our initially empty table:

	$o_1=x$	$o_2=z$	$o_3=y$
q^1			
q^2			

Filling the first column

	$o_1=x$	$o_2=z$	$o_3=y$
q^1	.6		
q^2	0		

$$v(1,1) = a_{<_s>1} \cdot b_1(x) = (1)(.6)$$

$$v(2,1) = a_{<_s>2} \cdot b_2(x) = (0)(.1)$$

Starting the second column

	$o_1=x$	$o_2=z$	$o_3=y$
q^1	.6		
q^2	0		

$$v(1,2) = \max_{i=1}^N v(i, 1) \cdot a_{i1} \cdot b_1(z)$$

$$= \max \begin{cases} v(1,1) \cdot a_{11} \cdot b_1(z) = (.6)(.7)(.3) \\ v(2,1) \cdot a_{21} \cdot b_1(z) = (0)(.5)(.3) \end{cases}$$

Starting the second column

	$o_1=x$	$o_2=z$	$o_3=y$
q^1	.6	.126	
q^2	0		

$$v(1,2) = \max_{i=1}^N v(i, 1) \cdot a_{i1} \cdot b_1(z)$$

$$= \max \begin{cases} v(1,1) \cdot a_{11} \cdot b_1(z) = (.6)(.7)(.3) \\ v(2,1) \cdot a_{21} \cdot b_1(z) = (0)(.5)(.3) \end{cases} \leftarrow$$

Finishing the second column

	$o_1=x$	$o_2=z$	$o_3=y$
q^1	.6	.126	
q^2	0		

$$v(2,2) = \max_{i=1}^N v(i,1) \cdot a_{i2} \cdot b_2(z)$$

$$= \max \left\{ \begin{array}{l} v(1,1) \cdot a_{12} \cdot b_2(z) = (.6)(.3)(.2) \\ v(2,1) \cdot a_{22} \cdot b_2(z) = (0)(.5)(.2) \end{array} \right.$$

Finishing the second column

	$o_1=x$	$o_2=z$	$o_3=y$
q^1	.6	.126	
q^2	0	.036	

$$v(2,2) = \max_{i=1}^N v(i, 1) \cdot a_{i2} \cdot b_2(z)$$

$$= \max \begin{cases} v(1,1) \cdot a_{12} \cdot b_2(z) = (.6)(.3)(.2) \\ v(2,1) \cdot a_{22} \cdot b_2(z) = (0)(.5)(.2) \end{cases}$$

Third column

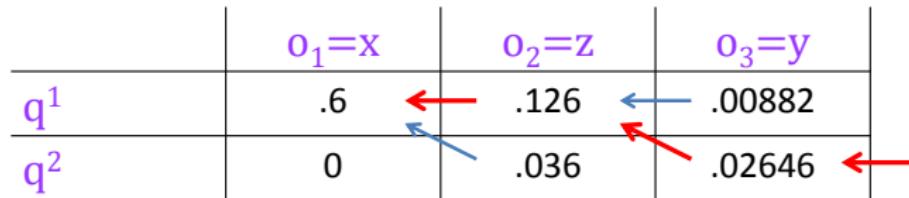
	$o_1=x$	$o_2=z$	$o_3=y$
q^1	.6	.126	.00882
q^2	0	.036	.02646

$$v(1, 3) = \max_{i=1}^N v(i, 2) a_{i1} b_1(y)$$

$$v(2, 3) = \max_{i=1}^N v(i, 2) a_{i2} b_2(y)$$

Best Path

	$o_1=x$	$o_2=z$	$o_3=y$
q^1	.6	.126	.00882
q^2	0	.036	.02646



- Choose best final state: $\max_{i=1}^N v(i, T)$
- Follow backtraces to find best full sequence: $q^1 q^1 q^2$

Viterbi pseudocode

```
function VITERBI(observations of len  $T$ ,state-graph of len  $N$ ) returns best-path
    create a path probability matrix  $viterbi[N+2,T]$ 
    for each state  $s$  from 1 to  $N$  do ; initialization step
         $viterbi[s,1] \leftarrow a_{0,s} * b_s(o_1)$ 
         $backpointer[s,1] \leftarrow 0$ 
    for each time step  $t$  from 2 to  $T$  do ; recursion step
        for each state  $s$  from 1 to  $N$  do
             $viterbi[s,t] \leftarrow \max_{s'=1}^N viterbi[s',t-1] * a_{s',s} * b_s(o_t)$ 
             $backpointer[s,t] \leftarrow \operatorname{argmax}_{s'=1}^N viterbi[s',t-1] * a_{s',s}$ 
     $viterbi[q_F,T] \leftarrow \max_{s=1}^N viterbi[s,T] * a_{s,q_F}$  ; termination step
     $backpointer[q_F,T] \leftarrow \operatorname{argmax}_{s=1}^N viterbi[s,T] * a_{s,q_F}$  ; termination step
    return the backtrace path by following backpointers to states back in time from
     $backpointer[q_F,T]$ 
```



What else?

- find the best tags for a sentence (**decoding**), and get $P(O, Q|\lambda)$ (**Viterbi algorithm**)

What else?

- find the best tags for a sentence (**decoding**), and get $P(O, Q|\lambda)$ (**Viterbi algorithm**)
- Compute the likelihood $P(O|\lambda)$, i.e., the probability of a sentence regardless of tags (a language model) (**Forward algorithm**)

$$\alpha(j, t) = \sum_{i=1}^N \alpha(i, t-1) \cdot a_{ij} \cdot b_j(o_t)$$

What else?

- find the best tags for a sentence (**decoding**), and get $P(O, Q|\lambda)$ (**Viterbi algorithm**)
- Compute the likelihood $P(O|\lambda)$, i.e., the probability of a sentence regardless of tags (a language model) (**Forward algorithm**)

$$\alpha(j, t) = \sum_{i=1}^N \alpha(i, t-1) \cdot a_{ij} \cdot b_j(o_t)$$

- Learn the best set of parameters $\lambda = (A, B)$, given only an *unannotated* corpus of sentences (unsupervised estimation) (**Forward-backward algorithm**) (not covered)

Maximum Entropy Markov Model (MEMM)

- HMMs : compute likelihood (observation word conditioned on tags)

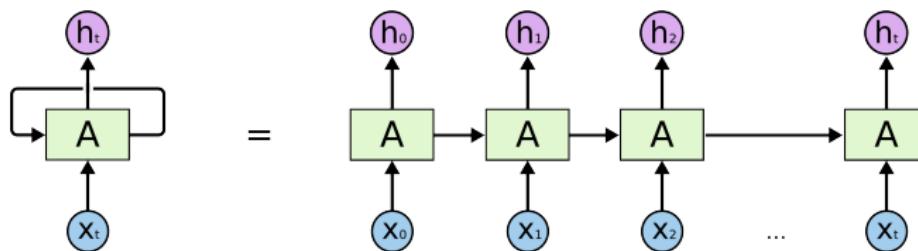
$$\hat{t}_1^n = \arg \max_{t_1^n} \prod_i P(w_i|t_i)P(t_i|t_{i-1})$$

- MEMM : Directly compute the posterior $P(t_1^n|w_1^n)$ (tags conditioned on observation words)

$$\hat{t}_1^n = \arg \max_{t_1^n} \prod_i P(t_i|w_i, t_{i-1})$$

Recurrent Neural Networks (RNNs)

- RNNs are used to model sequences : they are extremely good at language modelling and sequence labelling
- Traditional neural networks do not have *persistence*: presented with a new input, they forget the old one; RNNs solve this problem by having loops



- Input to the hidden layer is augmented with the activation value of the hidden layer from a previous point in time

Computation at the hidden layer

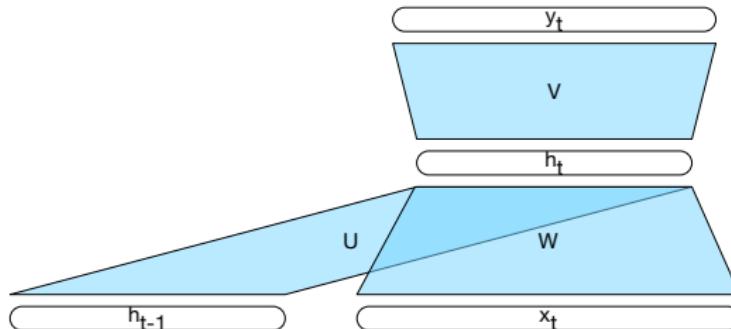
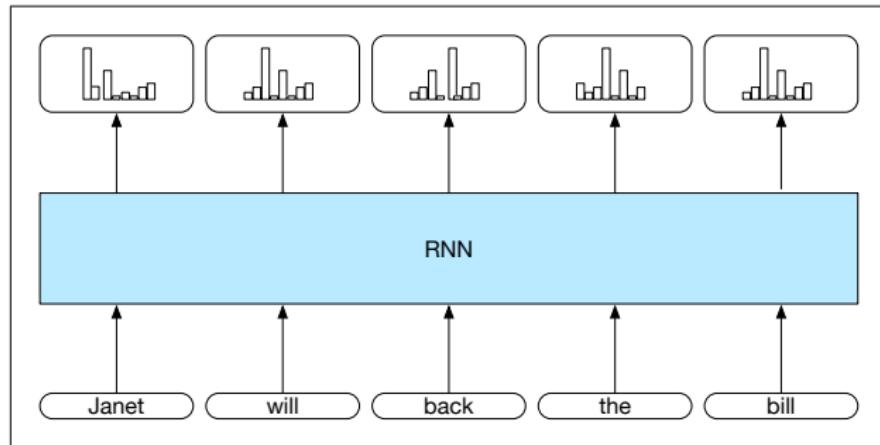


Figure 9.3 Simple recurrent neural network illustrated as a feed-forward network.

- A new set of weights U connect the hidden layer from the previous timestep to the current hidden layer
- U determines how the network makes use of past context in calculating the current output
- Just like other weights, these are also trained with backpropagation

$$h_t = g(Uh_{t-1} + Wx_t)$$

POS tagging as sequence labelling with RNN



- inputs are pre-trained word embeddings
- RNN block: input, hidden, output layer + U, W, V weight matrices
- output at each time step is distribution over POS tagset generated by softmax layer

Sequence labelling with RNNs and variants

- sequence of probability distributions (from softmax) can be combined with a tag-level language model, and selecting the most likely tag sequence
- Bi-directional RNNs, LSTMs, bi-LSTMs
- Sequence-to-sequence models, Transformers.

-
-

-
-

Sequence modelling

- Language Modelling: Sentence probability $P(w_1 \dots w_n)$; Guess the next word: $P(w_i | w_1 \dots w_{i-1})$
- POS tagging
- In general, these are insufficient models of language, because language has long distance dependencies

The train that my friends are arriving on from Paris tomorrow is cancelled.

Long distance dependencies in language

- Syntactic dependencies (Number/Gender agreement):

The **boy** in the house **was** watching television.

The **boys** in the house **were** watching television

She is planning to go to the conference **herself**

He is planning to go to the conference **himself**

Long distance dependencies in language

- Syntactic dependencies (Number/Gender agreement):

The **boy** in the house **was** watching television.

The **boys** in the house **were** watching television

She is planning to go to the conference **herself**

He is planning to go to the conference **himself**

- Semantic dependencies:

The **bird** next to the large oak tree near the grocery store on the corner **flies** rapidly.

The **man** next to the large oak tree near the grocery store on the corner **talks** rapidly.

Supertagging

- Like POS tagging, but tags are more complex, and more numerous
- Instead of 45-200 POS tags, 1000-2000 **supertags**
- Supertags capture syntactic information
- Supertags can be based on one of many "strongly lexicalised" grammar formalisms
 - * Combinatory Categorial Grammar (Steedman, 1996, 2000)
 - * Type-logical Grammars (Moortgat & Moot, 2002)
 - * Tree-adjoining grammars (Bangalore & Joshi 1999)

Combinatory Categorial Grammar (CCG)

- Categorial grammar CG is one of the oldest grammar formalisms (Ajdukiewicz, 1935; Bar-Hillel, 1953; Lambek 1958)
- Various flavours of CG now available: type-logical CG, algebraic pre-group grammars (Lambek), CCG
- CCG is now an established linguistic formalism (Steedman, 1996, 2000)
 - * syntax; semantics; prosody; information structure; wide-coverage parsing; incremental parsing, generation, etc.

Combinatory Categorial Grammar (CCG)

- CCG is a **strongly lexicalised** grammar

Combinatory Categorial Grammar (CCG)

- CCG is a **strongly lexicalised** grammar
- An elementary syntactic structure - a **lexical category** - is assigned to each word in a sentence.

walked : $S \setminus NP$ (*give me an NP to my left and I return a sentence*)

Combinatory Categorial Grammar (CCG)

- CCG is a **strongly lexicalised** grammar
- An elementary syntactic structure - a **lexical category** - is assigned to each word in a sentence.

walked : $S \setminus NP$ (*give me an NP to my left and I return a sentence*)

- A semantic interpretation is also assigned to each word

walked : $\lambda x.WALK(x)$

- A small number of rules define how categories can combine - rules are called **combinators**

CCG Lexical Categories

- Atomic categories : S , N , NP , PP (not many more)
(sentence, noun, noun phrase, prepositional phrase)
- Complex categories are built recursively from atomic categories and slashes, which indicate the directions of arguments
- Complex categories encode 'subcategorisation' information
 - * intransitive verb: $S \setminus NP$: *walked*
 - * transitive verb: $(S \setminus NP)/NP$: *respected*
 - * ditransitive verb : $((S \setminus NP)/NP)/NP$: *gave*
 - * np-pp : $((S \setminus NP)/PP)/NP$: *put*

A Simple CCG Derivation

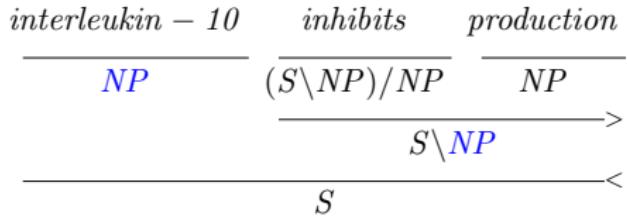
$$\frac{\text{interleukin} - 10 \quad \text{inhibits} \quad \text{production}}{NP \qquad \qquad \overbrace{(S \setminus NP) / NP \quad \overbrace{NP}^{\text{NP}}}^{(S \setminus NP)}} \rightarrow S \setminus NP$$

A Simple CCG Derivation

$$\frac{\begin{array}{ccc} \text{interleukin - 10} & \text{inhibits} & \text{production} \\ \hline NP & (S \setminus NP) / \textcolor{blue}{NP} & \textcolor{blue}{NP} \end{array}}{S \setminus NP}$$

> forward application

A Simple CCG Derivation



- > forward application
- < backward application

Function Application Rule Schemata

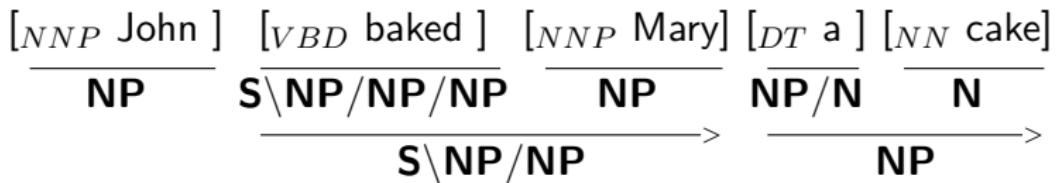
- Forward ($>$) and backward ($<$) application:

$$\begin{array}{lcl} X/Y \quad Y & \Rightarrow & X \quad (>) \\ Y \quad X\backslash Y & \Rightarrow & X \quad (<) \end{array}$$

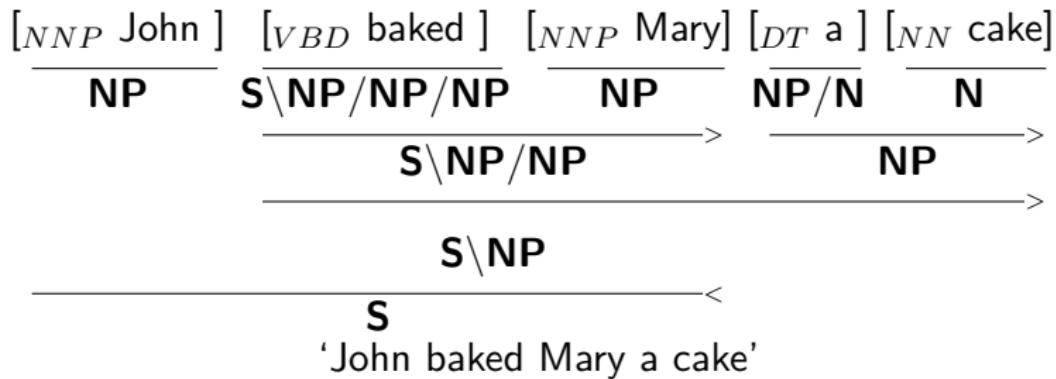
A Simple ccg Derivation

[*NNP* John] [*VBD* baked] [*NNP* Mary] [*DT* a] [*NN* cake]
NP S\NP/NP/NP NP NP/N N

A Simple ccg Derivation



A Simple ccg Derivation



baked : ((S \ NP) / NP) / NP

$\lambda x. \lambda y. \lambda z. z$ baked y for x

Syntactic combinatory potential for every word

bake :

(S \ NP) / NP / NP
(S \ NP) / PP / NP
(S \ NP) / NP
(S \ NP) / NP
(S \ NP) / S [to]
(S \ NP) / PP
.....

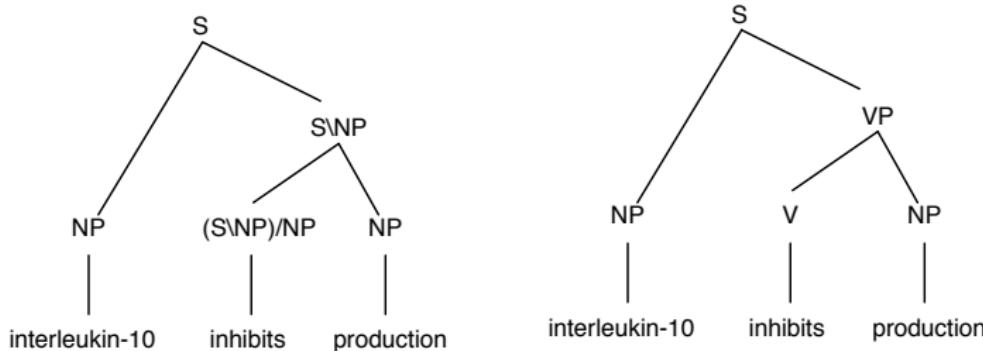
Statistical model

baked Mary a cake
baked a cake for Mary
baked a cake
baked to cheer himself up
baked for a charity

CCG Bank (English): > 1300 categories, 566 for verbs

Classical Categorial Grammar

- Only has application, and is context-free
- CCG has more combinators that make it "mildly" context sensitive.



Extraction out of a Relative Clause

The company which Microsoft bought

$$\frac{NP/N}{\overline{NP/N}} \quad \frac{N}{\overline{N}} \quad \frac{(NP\backslash NP)/(S/NP)}{\overline{(NP\backslash NP)/(S/NP)}} \quad \frac{NP}{\overline{NP}} \quad \frac{(S\backslash NP)/NP}{\overline{(S\backslash NP)/NP}}$$

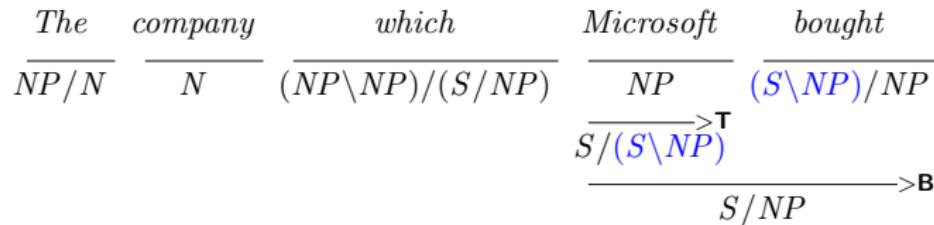
Extraction out of a Relative Clause

The company which Microsoft bought

$$\frac{NP/N}{N} \quad \frac{N}{(NP \setminus NP)/(S/NP)} \quad \frac{(NP \setminus NP)/(S/NP)}{NP} \quad \frac{NP}{(S \setminus NP)/NP}$$
$$S/(S \setminus NP) \xrightarrow{\mathbf{T}}$$

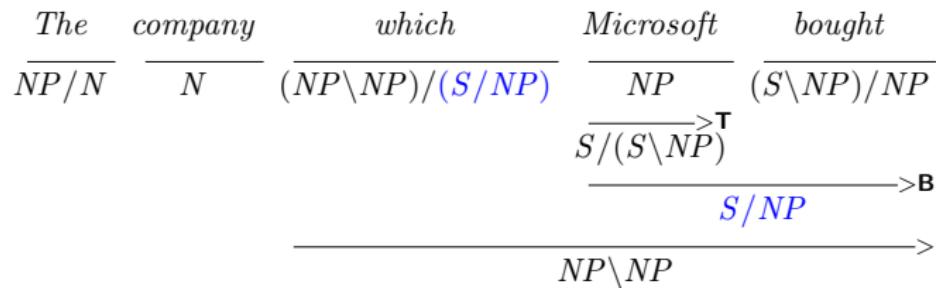
> Type-raising

Extraction out of a Relative Clause

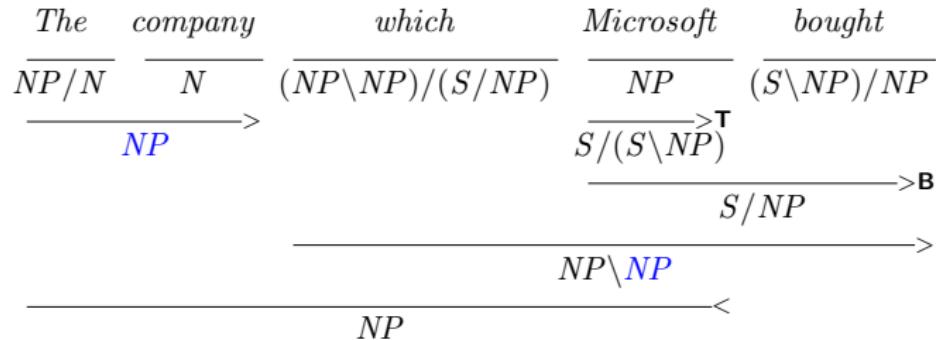


- > **T** Type-raising
- > **B** forward composition

Extraction out of a Relative Clause



Extraction out of a Relative Clause



Forward Composition and Type-Raising

- Forward composition ($>_B$):

$$X/Y \quad Y/Z \Rightarrow X/Z \quad (>_B)$$

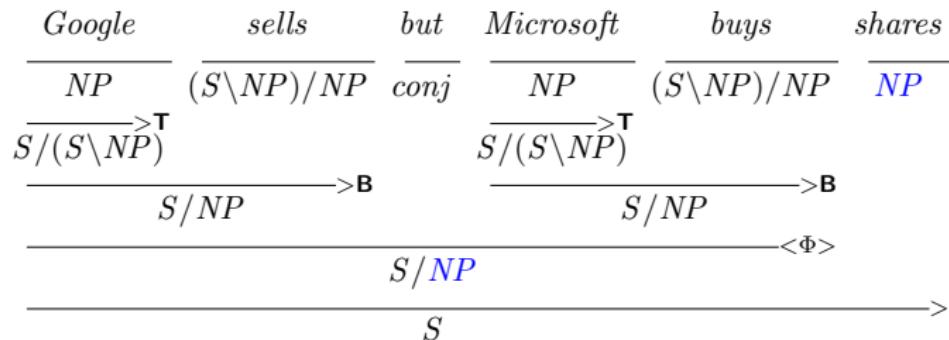
- Type-raising (\mathbf{T}):

$$X \Rightarrow T/(T \setminus X) \quad (>_T)$$

$$X \Rightarrow T \setminus (T/X) \quad (<_T)$$

- Extra combinatory rules increase the weak generative power to mild context-sensitivity

Non-constituents in CCG (Right Node Raising)



- > T type-raising
- > B forward composition

Why CCG/other categorial/tag grammars

- Formally, generative capacity matches human languages: mild-context sensitivity
- Transparent (one-to-one) mapping between syntax and semantics
- CCG emerging as dominant framework for “deep” semantic tasks like question-answering, textual Entailment.
- Supertaggers (English) currently 95% accuracy, almost at par with POS taggers
 - * 1200 lexical category types in CCGBank (compare to 45 POS tags in PennTreeBank)