

# Pattern Recognition 2018

## Support Vector Machines

Ad Feelders

Universiteit Utrecht

# Support Vector Machines



# Overview

- ① Separable Case
- ② Kernel Functions
- ③ Allowing Errors (Soft Margin)
- ④ SVM's in R.

# Linear Classifier for two classes

Linear model

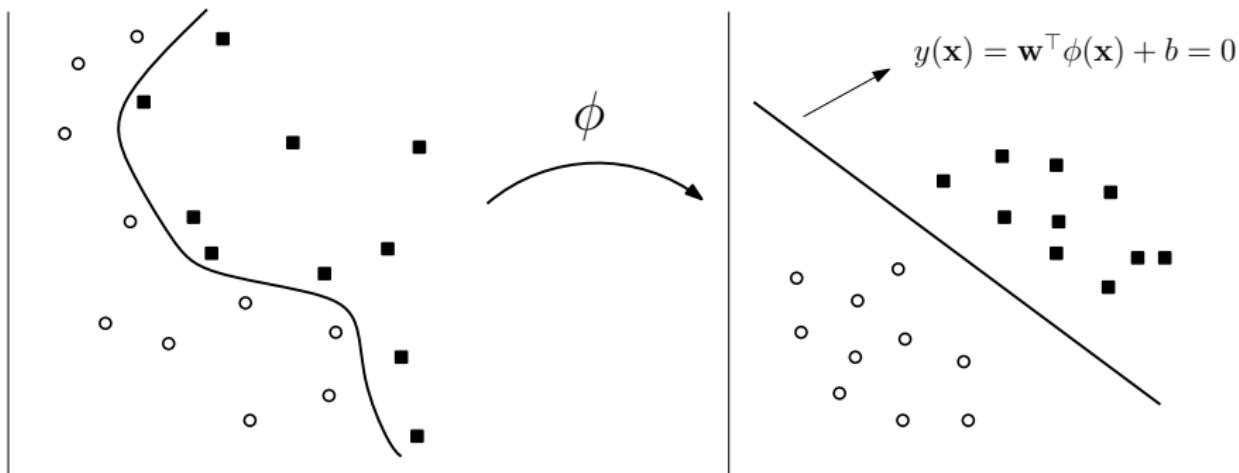
$$y(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) + b \quad (7.1)$$

with  $t_n \in \{-1, +1\}$ .

- Predict  $t_0 = +1$  if  $y(\mathbf{x}_0) \geq 0$  and  $t_0 = -1$  otherwise.
- The decision boundary is given by  $y(\mathbf{x}) = 0$ .

This is a linear classifier in feature space  $\phi(\mathbf{x})$ .

# Mapping



$\phi$  maps  $\mathbf{x}$  into higher dimensional space where data is linearly separable.

## Data linearly separable

Assume training data is linearly separable in feature space, so there is at least one choice of  $\mathbf{w}, b$  such that:

- ①  $y(\mathbf{x}_n) > 0$  for  $t_n = +1$ ;
- ②  $y(\mathbf{x}_n) < 0$  for  $t_n = -1$ ;

that is, all training points are classified correctly.

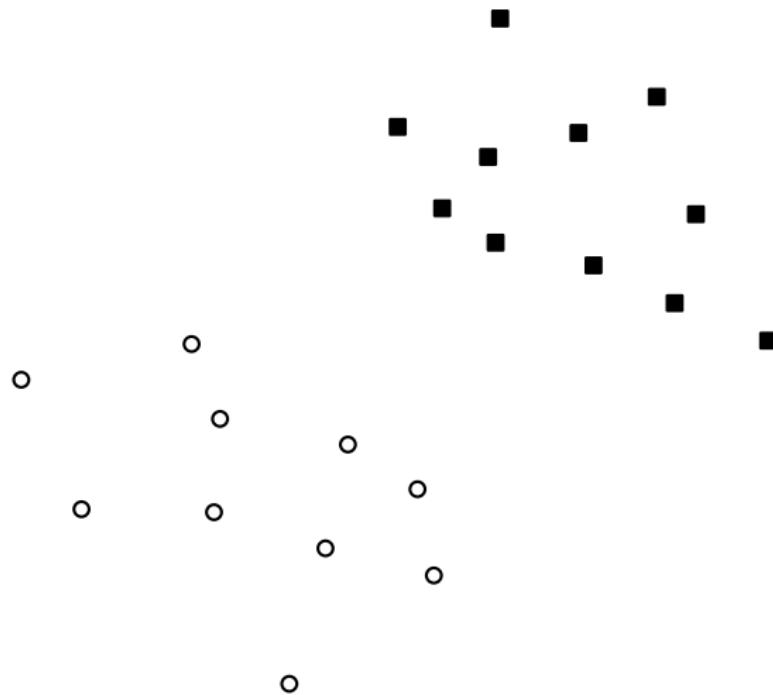
Putting 1. and 2. together:

$$t_n y(\mathbf{x}_n) > 0 \quad \text{for } n = 1, \dots, N$$

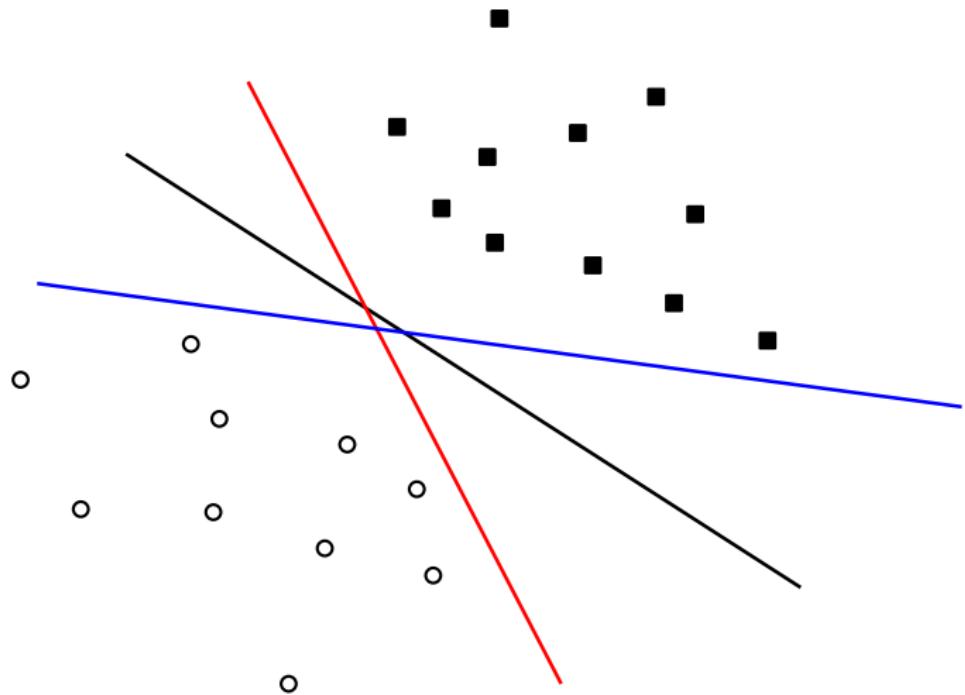
# Maximum Margin

- There may be many solutions that separate the classes exactly.
- Which one gives smallest prediction error?
- SVM chooses line with maximal *margin*, where the margin is the distance between the line and the closest data point.
- In this way, it avoids “low confidence” classifications.

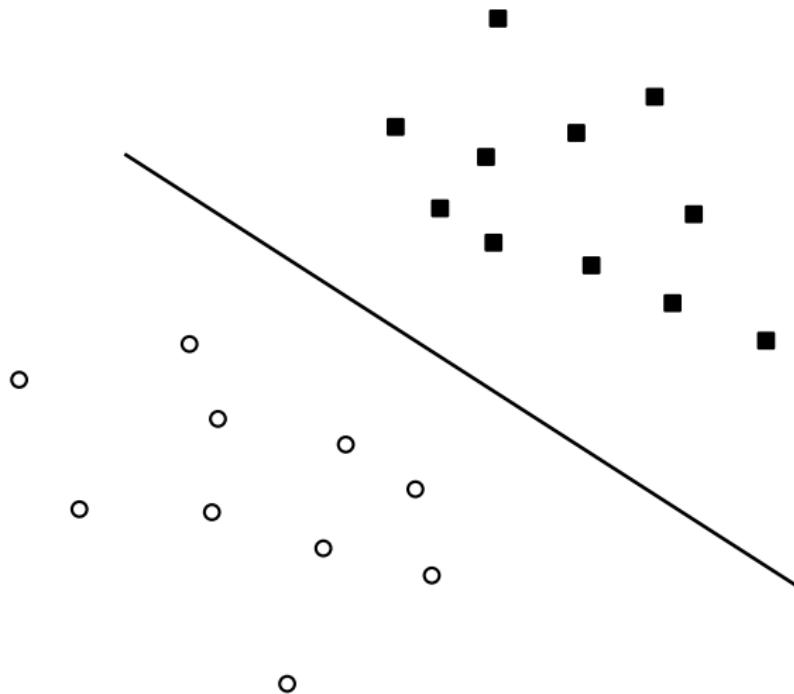
## Two-class training data



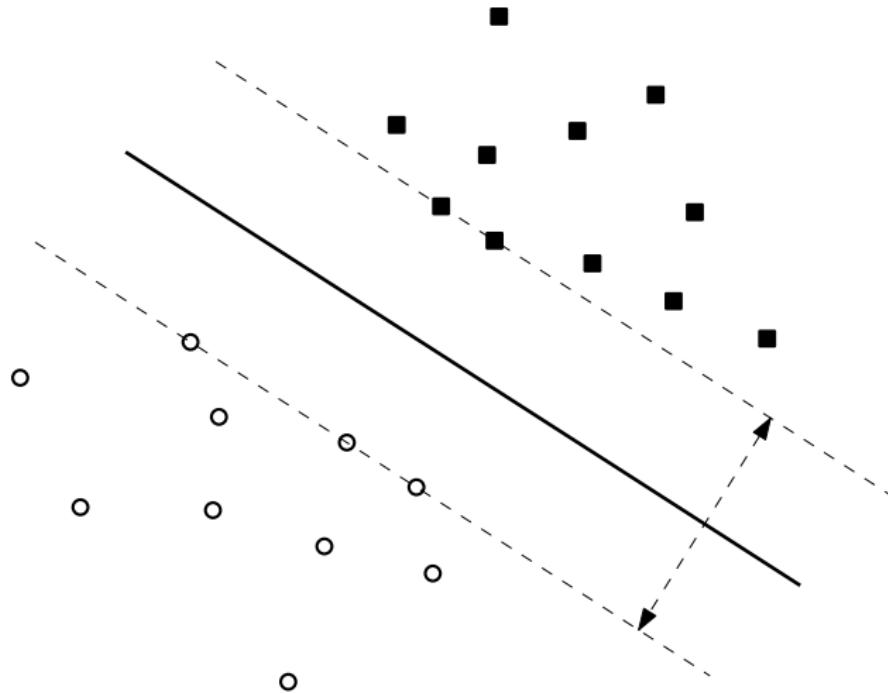
# Many Linear Separators



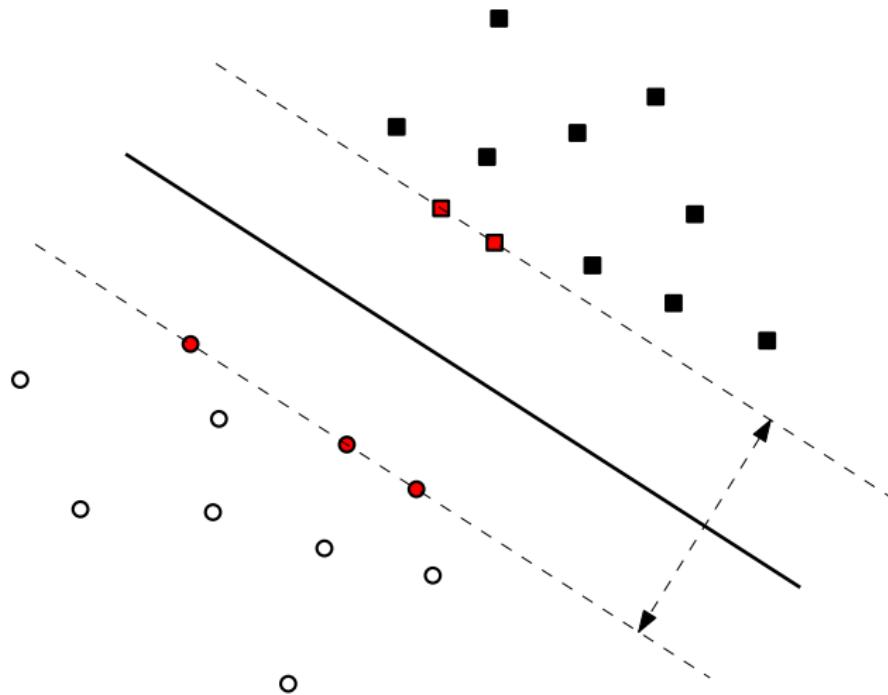
# SVM Decision Boundary



# Maximize Margin



# Support Vectors



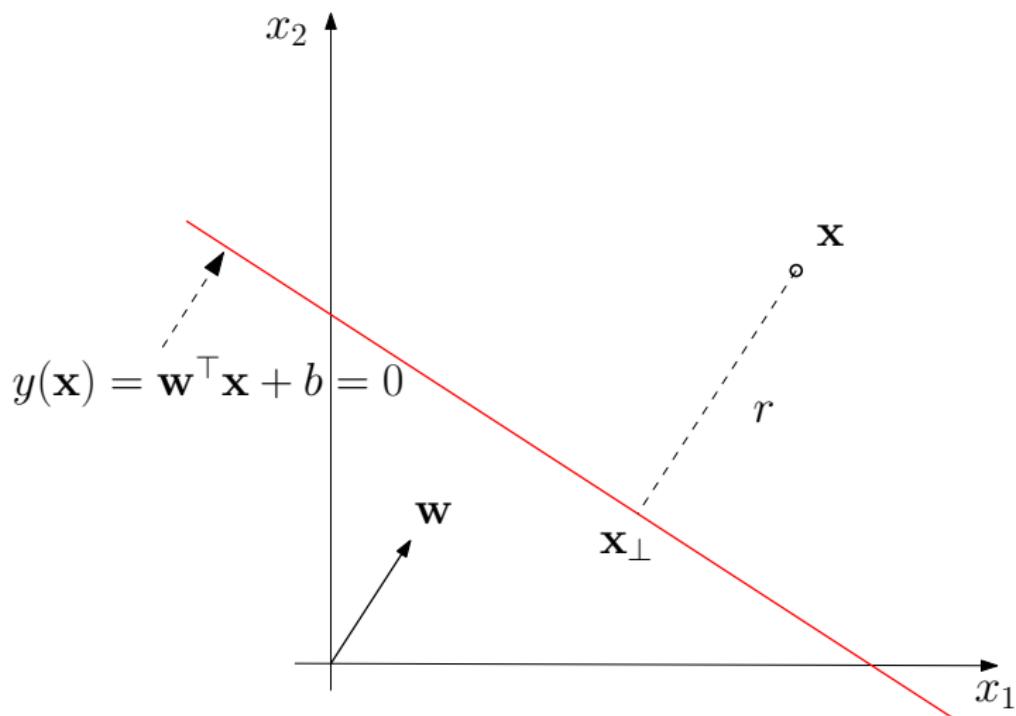
## Weight vector is orthogonal to the decision boundary

Consider two points  $\mathbf{x}_A$  and  $\mathbf{x}_B$  both of which lie on the decision surface. Because  $y(\mathbf{x}_A) = y(\mathbf{x}_B) = 0$ , we have

$$(\mathbf{w}^\top \mathbf{x}_A + b) - (\mathbf{w}^\top \mathbf{x}_B + b) = \mathbf{w}^\top (\mathbf{x}_A - \mathbf{x}_B) = 0$$

and so the vector  $\mathbf{w}$  is orthogonal to the decision surface.

# Distance of a point to a line



# Distance to decision surface ( $\phi(\mathbf{x}) = \mathbf{x}$ )

We have

$$\mathbf{x} = \mathbf{x}_\perp + r \frac{\mathbf{w}}{\|\mathbf{w}\|}. \quad (4.6)$$

where  $\frac{\mathbf{w}}{\|\mathbf{w}\|}$  is the unit vector in the direction of  $\mathbf{w}$ ,  $\mathbf{x}_\perp$  is the orthogonal projection of  $\mathbf{x}$  onto the line  $y(\mathbf{x}) = 0$ , and  $r$  is the (signed) distance of  $\mathbf{x}$  to the line. Multiply (4.6) left and right by  $\mathbf{w}^\top$  and add  $b$ :

$$\underbrace{\mathbf{w}^\top \mathbf{x} + b}_{y(\mathbf{x})} = \underbrace{\mathbf{w}^\top \mathbf{x}_\perp + b}_{0} + r \frac{\mathbf{w}^\top \mathbf{w}}{\|\mathbf{w}\|}$$

So we get

$$r = y(\mathbf{x}) \frac{\|\mathbf{w}\|}{\|\mathbf{w}\|^2} = \frac{y(\mathbf{x})}{\|\mathbf{w}\|} \quad (4.7)$$

## Distance of a point to a line

The signed distance of  $\mathbf{x}_n$  to the decision boundary is

$$r = \frac{y(\mathbf{x}_n)}{\|\mathbf{w}\|}$$

For lines that separate the data perfectly, we have  $t_n y(\mathbf{x}_n) = |y(\mathbf{x}_n)|$ , so that the distance is given by

$$\frac{t_n y(\mathbf{x}_n)}{\|\mathbf{w}\|} = \frac{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b)}{\|\mathbf{w}\|} \quad (7.2)$$

# Maximum margin solution

Now we are ready to define the optimization problem:

$$\arg \max_{\mathbf{w}, b} \left\{ \min_n \left[ \frac{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b)}{\|\mathbf{w}\|} \right] \right\}. \quad (7.3)$$

Since  $\frac{1}{\|\mathbf{w}\|}$  does not depend on  $n$ , it can be moved outside of the minimization:

$$\arg \max_{\mathbf{w}, b} \left\{ \frac{1}{\|\mathbf{w}\|} \min_n [t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b)] \right\}. \quad (7.3)$$

- Direct solution of this problem would be rather complex.
- A more convenient representation is possible.

# Canonical Representation

The hyperplane (decision boundary) is defined by

$$\mathbf{w}^\top \phi(\mathbf{x}) + b = 0$$

Then also

$$\kappa(\mathbf{w}^\top \phi(\mathbf{x}) + b) = \kappa\mathbf{w}^\top \phi(\mathbf{x}) + \kappa b = 0$$

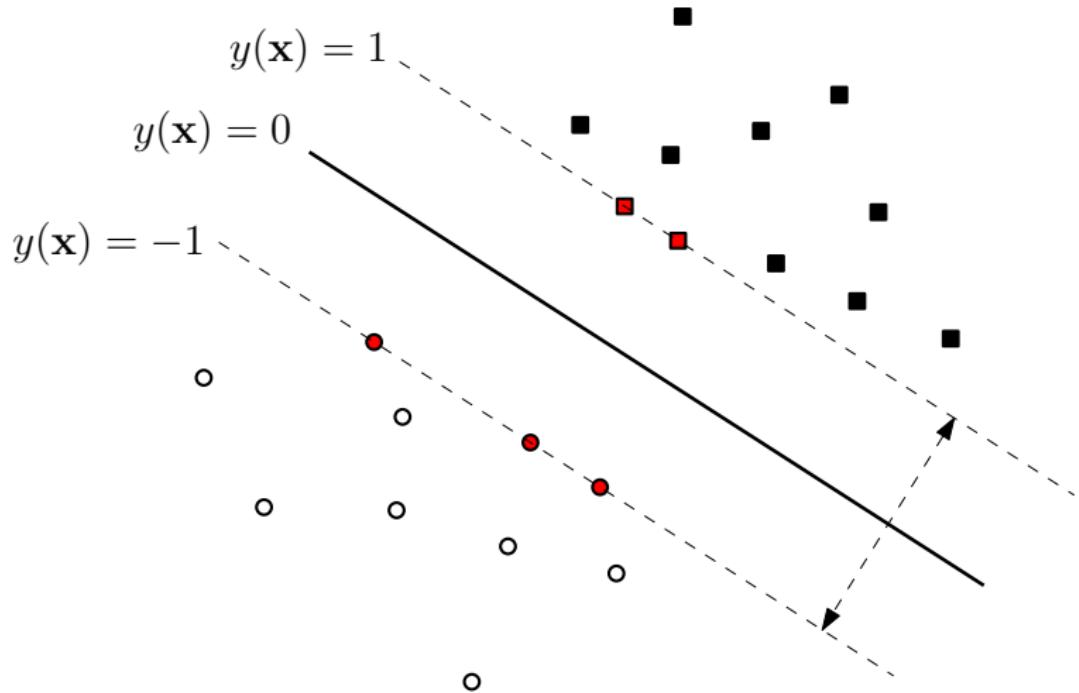
so rescaling  $\mathbf{w} \rightarrow \kappa\mathbf{w}$  and  $b \rightarrow \kappa b$  gives just another representation of the same decision boundary.

To resolve this ambiguity, we choose the scaling factor such that

$$t_i(\mathbf{w}^\top \phi(\mathbf{x}_i) + b) = 1 \tag{7.4}$$

for the points  $\mathbf{x}_i$  closest to the decision boundary.

## Canonical Representation (square=1,circle=-1)



# Canonical Representation

In this case we have

$$t_n(\mathbf{w}^\top \phi(\mathbf{x}_n) + b) \geq 1 \quad n = 1, \dots, N \quad (7.5)$$

Quadratic program

$$\arg \min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 \quad (7.6)$$

subject to the constraints (7.5).

This optimization problem has a unique global minimum.

# Lagrangian Function

Introduce Lagrange multipliers  $a_n \geq 0$  to get Lagrangian function

$$L(\mathbf{w}, b, \mathbf{a}) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N a_n \{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b) - 1\} \quad (7.7)$$

with

$$\frac{\partial L(\mathbf{w}, b, \mathbf{a})}{\partial \mathbf{w}} = \mathbf{w} - \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n)$$

# Lagrangian Function

and for  $b$ :

$$\frac{\partial L(\mathbf{w}, b, \mathbf{a})}{\partial b} = - \sum_{n=1}^N a_n t_n$$

Equating the derivatives to zero yields the conditions:

$$\mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n) \quad (7.8)$$

and

$$\sum_{n=1}^N a_n t_n = 0 \quad (7.9)$$

# Dual Representation

Eliminating  $\mathbf{w}$  and  $b$  from  $L(\mathbf{w}, b, \mathbf{a})$  gives the *dual representation*.

$$\begin{aligned} L(\mathbf{w}, b, \mathbf{a}) &= \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N a_n \{t_n (\mathbf{w}^\top \phi(\mathbf{x}_n) + b) - 1\} \\ &= \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N a_n t_n \mathbf{w}^\top \phi(\mathbf{x}_n) - b \sum_{n=1}^N a_n t_n + \sum_{n=1}^N a_n \\ &= \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m \phi(\mathbf{x}_n)^\top \phi(\mathbf{x}_m) \\ &\quad - \sum_{n=1}^N \sum_{m=1}^N a_n t_n a_m t_m \phi(\mathbf{x}_n)^\top \phi(\mathbf{x}_m) + \sum_{n=1}^N a_n \\ &= \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n t_n a_m t_m \phi(\mathbf{x}_n)^\top \phi(\mathbf{x}_m) \end{aligned}$$

# Dual Representation

Maximize

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n,m=1}^N a_n t_n a_m t_m \phi(\mathbf{x}_n)^\top \phi(\mathbf{x}_m) \quad (7.10)$$

with respect to  $\mathbf{a}$  and subject to the constraints

$$a_n \geq 0, \quad n = 1, \dots, N \quad (7.11)$$

$$\sum_{n=1}^N a_n t_n = 0. \quad (7.12)$$

# Kernel Function

- We map  $\mathbf{x}$  to a high-dimensional space  $\phi(\mathbf{x})$  in which data is linearly separable.
- Performing computations in this high-dimensional space may be very expensive.
- Use a kernel function  $k$  that computes a dot product in this space without making the actual mapping (“kernel trick”):

$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^\top \phi(\mathbf{x}')$$

## Example: polynomial kernel

Suppose  $\mathbf{x} \in \mathbb{R}^3$  and  $\phi(\mathbf{x}) \in \mathbb{R}^{10}$  with

$$\phi(\mathbf{x}) = (1, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_3, x_1^2, x_2^2, x_3^2, \sqrt{2}x_1x_2, \sqrt{2}x_1x_3, \sqrt{2}x_2x_3)$$

Then

$$\begin{aligned}\phi(\mathbf{x})^\top \phi(\mathbf{z}) &= 1 + 2x_1z_1 + 2x_2z_2 + 2x_3z_3 + x_1^2z_1^2 + x_2^2z_2^2 + x_3^2z_3^2 \\ &\quad + 2x_1x_2z_1z_2 + 2x_1x_3z_1z_3 + 2x_2x_3z_2z_3\end{aligned}$$

But this can be written as

$$(1 + \mathbf{x}^\top \mathbf{z})^2 = (1 + x_1z_1 + x_2z_2 + x_3z_3)^2$$

which costs much less operations to compute.

## Polynomial kernel: numeric example

Suppose  $\mathbf{x} = (3, 2, 6)$  and  $\mathbf{z} = (4, 1, 5)$ .

Then

$$\begin{aligned}\phi(\mathbf{x}) &= (1, 3\sqrt{2}, 2\sqrt{2}, 6\sqrt{2}, 9, 4, 36, 6\sqrt{2}, 18\sqrt{2}, 12\sqrt{2}) \\ \phi(\mathbf{z}) &= (1, 4\sqrt{2}, 1\sqrt{2}, 5\sqrt{2}, 16, 1, 25, 4\sqrt{2}, 20\sqrt{2}, 5\sqrt{2})\end{aligned}$$

Then

$$\phi(\mathbf{x})^\top \phi(\mathbf{z}) = 1 + 24 + 4 + 60 + 144 + 4 + 900 + 48 + 720 + 120 = 2025.$$

But

$$(1 + \mathbf{x}^\top \mathbf{z})^2 = (1 + (3)(4) + (2)(1) + (6)(5))^2 = 45^2 = 2025$$

is a more efficient way to compute this dot product.

# Kernels

Linear kernel

$$k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^\top \mathbf{x}'$$

Two popular non-linear kernels are the polynomial kernel (of degree  $M$ ):

$$k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}' + c)^M$$

and Gaussian (or radial) kernel:

$$k(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|^2 / 2\sigma^2), \quad (6.23)$$

or

$$k(\mathbf{x}, \mathbf{x}') = \exp(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2),$$

where  $\gamma = \frac{1}{2\sigma^2}$ .

## Dual Representation with kernels

Using  $k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^\top \phi(\mathbf{x}')$  we get dual representation:

Maximize

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n,m=1}^N a_n t_n a_m t_m k(\mathbf{x}_n, \mathbf{x}_m) \quad (7.10)$$

with respect to  $\mathbf{a}$  and subject to the constraints

$$a_n \geq 0, \quad n = 1, \dots, N \quad (7.11)$$

$$\sum_{n=1}^N a_n t_n = 0. \quad (7.12)$$

Is this dual “easier” than the original problem?

# Prediction

Recall that

$$y(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) + b \quad (7.1)$$

Substituting

$$\mathbf{w} = \sum_{n=1}^N a_n t_n \phi(\mathbf{x}_n) \quad (7.8)$$

into (7.1), we get

$$y(\mathbf{x}) = b + \sum_{n=1}^N a_n t_n k(\mathbf{x}, \mathbf{x}_n) \quad (7.13)$$

# Constrained Optimization

Minimize

$$f(x)$$

subject to

$$g_i(x) \geq 0$$

Lagrangian function:

$$L(x, \lambda) = f(x) - \sum_i \lambda_i g_i(x)$$

KKT conditions for solution:

- ①  $g_i(x) \geq 0$
- ②  $\lambda_i \geq 0$
- ③  $\lambda_i g_i(x) = 0$

## Prediction: support vectors

KKT conditions:

$$a_n \geq 0 \quad (7.14)$$

$$t_n y(\mathbf{x}_n) - 1 \geq 0 \quad (7.15)$$

$$a_n \{ t_n y(\mathbf{x}_n) - 1 \} = 0 \quad (7.16)$$

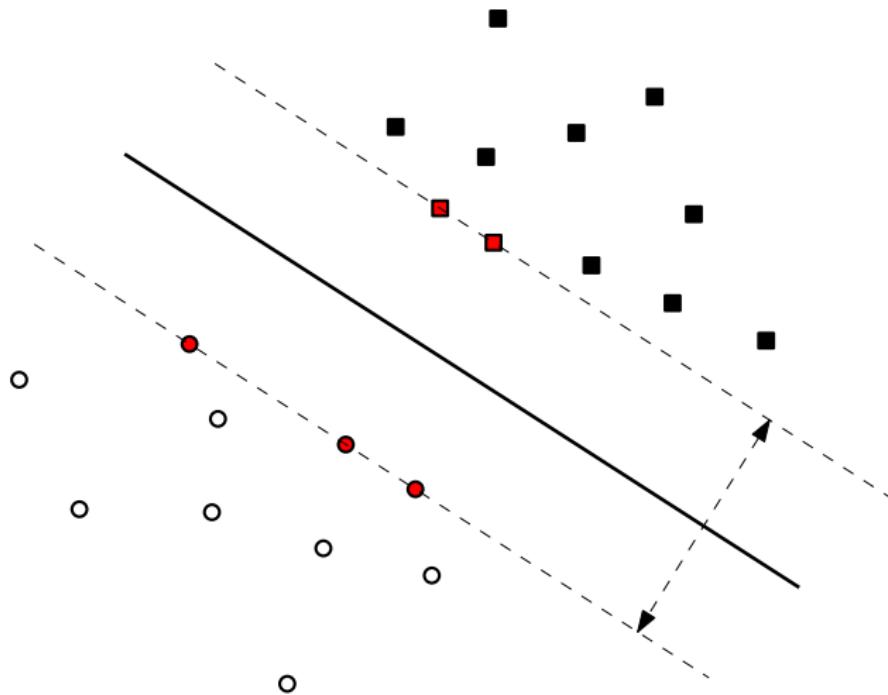
From (7.16) it follows that for every data point, either

- ①  $a_n = 0$ , or
- ②  $t_n y(\mathbf{x}_n) = 1$ .

The former play no role in making predictions (see 7.13), and the latter are the *support vectors* that lie on the maximum margin hyper planes.

Only the support vectors play a role in predicting the class of new attribute vectors!

# Only the support vectors are important for prediction



## Prediction: computing $b$

Since for any support vector  $\mathbf{x}_n$  we have  $t_n y(\mathbf{x}_n) = 1$ , we can use (7.13) to get

$$t_n \left( b + \sum_{m \in \mathcal{S}} a_m t_m k(\mathbf{x}_n, \mathbf{x}_m) \right) = 1, \quad (7.17)$$

where  $\mathcal{S}$  denotes the set of support vectors.

Hence we have

$$t_n b + t_n \sum_{m \in \mathcal{S}} a_m t_m k(\mathbf{x}_n, \mathbf{x}_m) = 1$$

$$t_n b = 1 - t_n \sum_{m \in \mathcal{S}} a_m t_m k(\mathbf{x}_n, \mathbf{x}_m)$$

and since  $t_n \in \{-1, +1\}$  and so  $1/t_n = t_n$ :

$$b = t_n - \sum_{m \in \mathcal{S}} a_m t_m k(\mathbf{x}_n, \mathbf{x}_m) \quad (7.17a)$$

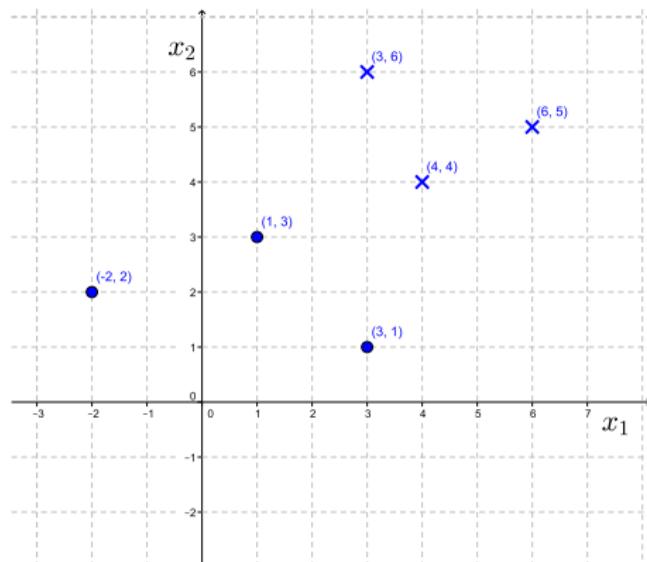
## Prediction: Example

We receive the following output from the optimization software for fitting a support vector machine with linear kernel and perfect separation of the training data:

$n$	$x_{n,1}$	$x_{n,2}$	$t_n$	$a_n$
1	-2	2	-1	0
2	1	3	-1	$\frac{1}{8}$
3	3	1	-1	$\frac{1}{8}$
4	3	6	+1	0
5	4	4	+1	$\frac{1}{4}$
6	6	5	+1	0

## Prediction: Example

The figure below is a plot of the same data set, where the dots represent points with class  $-1$ , and the crosses points with class  $+1$ .



## Prediction: Example

- (a) Compute the value of the SVM bias term  $b$ .

Data points with  $a > 0$  are support vectors.

Let's take the point  $x_1 = 4, x_2 = 4$  with class label +1:

$$b = t_m - \sum_{n=1}^N a_n t_n \mathbf{x}_m^\top \mathbf{x}_n = 1 + \frac{1}{8}[4 \ 4] \begin{bmatrix} 1 \\ 3 \end{bmatrix} + \frac{1}{8}[4 \ 4] \begin{bmatrix} 3 \\ 1 \end{bmatrix} - \frac{1}{4}[4 \ 4] \begin{bmatrix} 4 \\ 4 \end{bmatrix} = -3$$

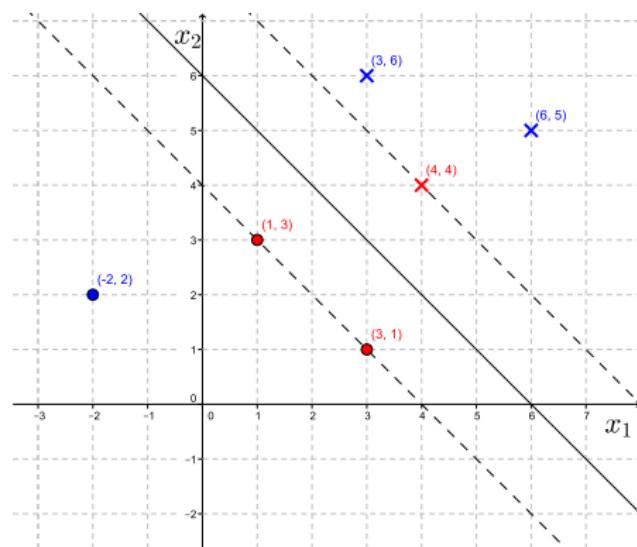
- (b) Which class does the SVM predict for the data point  $x_1 = 5, x_2 = 2$ ?

$$y(\mathbf{x}) = b + \sum_{n=1}^N a_n t_n \mathbf{x}^\top \mathbf{x}_n = -3 - \frac{1}{8}[5 \ 2] \begin{bmatrix} 1 \\ 3 \end{bmatrix} - \frac{1}{8}[5 \ 2] \begin{bmatrix} 3 \\ 1 \end{bmatrix} + \frac{1}{4}[5 \ 2] \begin{bmatrix} 4 \\ 4 \end{bmatrix} = \frac{1}{2}$$

Since the sign is positive, we predict class +1.

# Prediction: Example

Decision boundary and support vectors.



# Allowing Errors

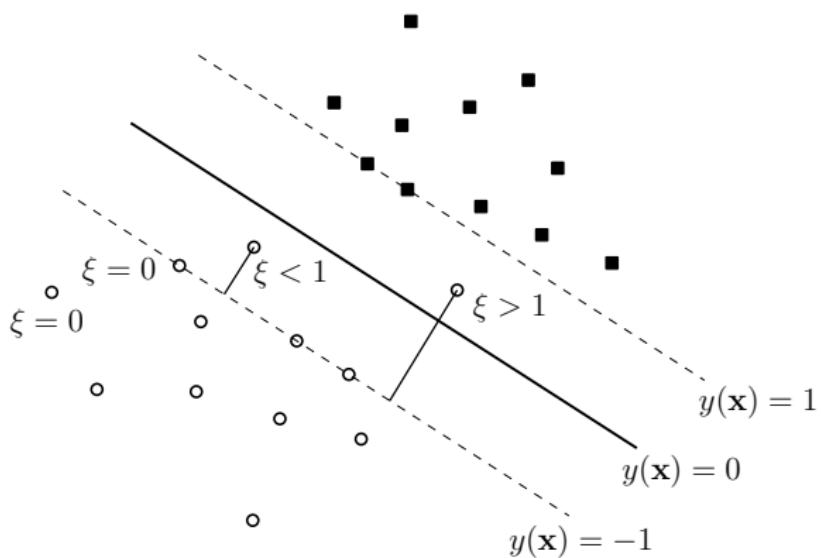
- So far we assumed that the training data points are linearly separable in feature space  $\phi(\mathbf{x})$ .
- Resulting SVM gives exact separation of training data in original input space  $\mathbf{x}$ , with non-linear decision boundary.
- Class distributions typically overlap, in which case exact separation of the training data leads to poor generalization (overfitting).

# Allowing Errors

- Data points are allowed to be on the “wrong” side of the margin boundary, but with a penalty that increases with the distance from that boundary.
- For convenience we make this penalty a linear function of the distance to the margin boundary.
- Introduce slack variables  $\xi_n \geq 0$  with one slack variable for each training data point.

# Definition of Slack Variables

We define  $\xi_n = 0$  for data points that are on the inside of the correct margin boundary and  $\xi_n = |t_n - y(\mathbf{x}_n)|$  for all other data points.



## New objective function

Our goal is to maximize the margin while softly penalizing points that lie on the wrong side of the margin boundary. We therefore minimize

$$C \sum_{n=1}^N \xi_n + \frac{1}{2} \|\mathbf{w}\|^2 \quad (7.21)$$

where the parameter  $C > 0$  controls the trade-off between the slack variable penalty and the margin. Alternative view (divide by  $C$  and put  $\lambda = \frac{1}{2C}$ ):

$$\sum_{n=1}^N \xi_n + \lambda \sum w_i^2$$

First term represents lack-of-fit (hinge loss) and second term takes care of regularization.

# Dual

Maximize

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n,m=1}^N a_n t_n a_m t_m k(\mathbf{x}_n, \mathbf{x}_m) \quad (7.32)$$

with respect to  $\mathbf{a}$  and subject to the constraints

$$0 \leq a_n \leq C, \quad n = 1, \dots, N \quad (7.33)$$

$$\sum_{n=1}^N a_n t_n = 0. \quad (7.34)$$

The same as the separable case, except for the constraints  $a_n \leq C$ .

# Model Selection

- As usual we are confronted with the problem of selecting the appropriate model complexity.
- The relevant parameters are  $C$  and any parameters of the chosen kernel function.

# SVM in R

LIBSVM is available in package e1071 in R.

It can also perform regression and non-binary classification.

Non-binary classification is performed as follows:

- Train  $K(K - 1)/2$  binary SVM's on all possible pairs of classes.
- To classify a new point, let it be classified by every binary SVM, and pick the class with the highest number of votes.
- This is done automatically by function `svm` in e1071.

# How to in R: analysis of optdigits data

```
# SVM with radial kernel and gamma=1/62 and cost=1 (default settings)

> optdigits.svm <- svm(optdigits.train[, -c(1, 40,65)],optdigits.train[,65])

# make predictions on test set

> svm.pred <- predict(optdigits.svm,optdigits.test[, -c(1, 40,65)])

> table(optdigits.test[,65],svm.pred)

  svm.pred
    0   1   2   3   4   5   6   7   8   9
0 177   0   0   0   1   0   0   0   0   0
1   0 179   0   0   2   0   0   0   1   0
2   0   7 167   0   3   0   0   0   0   0
3   0   0   3 172   2   2   0   1   2   1
4   0   1   0   0 179   0   0   0   1   0
5   0   0   0   0   0 181   0   0   0   1
6   1   0   0   0   1   0 179   0   0   0
7   0   0   0   0   0   0   0 172   0   7
8   0   6   0   0   3   0   0   0 159   6
9   0   0   0   2   0   1   1   1   1 174

# accuracy on test sample is somewhat better than regularized multinomial logit
# which had "only" 95% accuracy

> sum(diag(table(optdigits.test[,65],svm.pred)))/nrow(optdigits.test)
[1] 0.967724
```

# How to in R: : analysis of optdigits data

```
# tune cost parameter with cross-validation
> optdigits.svm.tune <- tune.svm(optdigits.train[, -c(1, 40,65)],
                                    optdigits.train[,65],cost=1:10)

Warning messages:
1: In svm.default(list(V2 = c(0L, 0L, 6L, 0L, 3L, 1L, 0L, 0L, 0L, 0L,  :
   Variable(s) V57 constant. Cannot scale data.
# V57 is almost always zero, let's remove it
> optdigits.svm.tune <- tune.svm(optdigits.train[, -c(1, 40,57,65)],
                                    optdigits.train[,65],cost=1:10)

# show performance for each value of the cost parameter
> optdigits.svm.tune$performances
  cost      error dispersion
1    1 0.01490643 0.005227509
2    2 0.01202958 0.005934869
3    3 0.01229136 0.005376797
4    4 0.01176917 0.005119188
5    5 0.01150739 0.004803942
6    6 0.01150739 0.004803942
7    7 0.01176849 0.005115891
8    8 0.01150739 0.004803942
9    9 0.01176849 0.005115891
10   10 0.01176849 0.005115891
```

## How to in R: : analysis of optdigits data

```
# cost=5 is the smallest among the parameter values that minimize  
# cross-validation error  
# We fit the SVM with this parameter value on the whole training set:  
  
> optdigits.svm.tuned <- svm(optdigits.train[, -c(1, 40,57,65)],  
                           optdigits.train[,65],cost=5)  
  
# Generate predictions on the test set  
  
> svm.tuned.pred <- predict(optdigits.svm.tuned,  
                           optdigits.test[, -c(1, 40,57,65)])  
  
# Compute accuracy on test set  
  
> sum(diag(table(optdigits.test[,65],svm.tuned.pred)))/nrow(optdigits.test)  
[1] 0.9727323
```